

Models of retrieval in sentence comprehension: A computational evaluation using Bayesian
hierarchical modeling

Bruno Nicenboim

Department of Linguistics, University of Potsdam,
Potsdam, Germany

E-mail:bruno.nicenboim@uni-potsdam.de

Shravan Vasishth

Department of Linguistics, University of Potsdam, Potsdam, Germany

Draft of June 25, 2017.

Author Note

We thank Hedderik van Rijn for the useful discussion during Groningen Spring School on Cognitive Modeling, 2016, and Pavel Logačev for insightful comments on the manuscript. This work was supported by Potsdam Graduate School, the University of Potsdam, and by Volkswagen Foundation grant 89953.

Abstract

Research on similarity-based interference has provided extensive evidence that the formation of dependencies between non-adjacent words relies on a cue-based retrieval mechanism. There are two different models that can account for one of the main predictions of interference, i.e., a slowdown at a retrieval site, when several items share a feature associated with a retrieval cue: Lewis and Vasishth's (2005) activation-based model and McElree's (2000) direct-access model. Even though these two models have been used almost interchangeably, they are based on different assumptions and predict differences in the relationship between reading times and response accuracy. The activation-based model follows the assumptions of the ACT-R framework, and its retrieval process behaves as a lognormal race between accumulators of evidence with a single variance. Under this model, accuracy of the retrieval is determined by the winner of the race and retrieval time by its rate of accumulation. In contrast, the direct-access model assumes a model of memory where only the probability of retrieval can be affected, while the retrieval time is drawn from the same distribution; in this model, differences in latencies are a by-product of the possibility of backtracking and repairing incorrect retrievals. We implemented both models in a Bayesian hierarchical framework in order to evaluate them and compare them. The data show that correct retrievals take longer than incorrect ones, and this pattern is better fit under the direct-access model than under the activation-based model. This finding does not rule out the possibility that retrieval may be behaving as a race model with assumptions that follow less closely the ones from the ACT-R framework. By introducing a modification of the activation model, i.e, by assuming that the accumulation of evidence for retrieval of incorrect items is not only slower but noisier (i.e., different variances for the correct and incorrect items), the model can provide a fit as good as the one of the direct-access model. This first ever computational evaluation of alternative accounts of retrieval processes in sentence processing opens the way for a broader investigation of theories of dependency completion.

Keywords: cognitive modeling; sentence processing; working memory; cue-based retrieval; similarity-based interference; Bayesian hierarchical modeling

Models of retrieval in sentence comprehension: A computational evaluation using Bayesian hierarchical modeling

There is strong evidence that the formation of syntactic dependencies between non-adjacent words relies on the memory system. An example is the so-called locality effect: increasing the distance between co-dependents (such as subjects and verbs) tends to lead to greater processing difficulty (Gibson, 2000; Grodner & Gibson, 2005). Research on interference makes a similar point: the speed and/or accuracy of dependency completion can be adversely affected by the presence of items in memory that are similar to the retrieval target (among others: Gordon, Hendrick, & Levine, 2002; Van Dyke & Lewis, 2003; Van Dyke, 2007; Jäger, Engelmann, & Vasishth, 2015; Nicenboim, Engelmann, Suckow, & Vasishth, Submitted; Vasishth, Brüßow, Lewis, & Drenhaus, 2008). Such a central role for memory in sentence comprehension is well-motivated: it is implausible that the parser could keep track of a large and in principle unbounded inventory of the dependencies that can be found in a sentence, since they easily exceed the amount of information that can be held in the focus of attention (McElree & Doshier, 1989; McElree, 2006; Cowan, 1995; Oberauer, 2013; Marcus, 2013). The evidence from studies investigating similarity-based interference (see the meta-analysis of published reading studies in Jäger, Engelmann, & Vasishth, 2017) suggests that dependency completion relies on a *content-addressable cue-based retrieval mechanism* that is subject to interference (McElree, 2000; Van Dyke & Lewis, 2003; Lewis, Vasishth, & Van Dyke, 2006). Similarity-based interference is a phenomenon that is not unique to language, and occurs when several items share a feature associated with a retrieval cue. A major implication is that the retrieval mechanism employed for the creation of linguistic dependencies is similar to the one utilized in non-language domains.

There are multiple implementations compatible with such a content-addressable cue-based retrieval mechanism in sentence processing. As a verbally stated model, this type of mechanism would entail that when retrieval cues fully match the target of retrieval,

similarity-based interference would cause an inhibitory effect, that is, an increase of processing difficulty at the retrieval of a dependent. This processing difficulty would be reflected in longer reading times and lower accuracy. However, in some cases, shorter reading times have been observed when increased processing difficulty was clearly expected (Van Dyke & McElree, 2006; Nicenboim, Vasishth, Gattei, Sigman, & Kliegl, 2015; Nicenboim, Logačev, Gattei, & Vasishth, 2016). In these cases, it is usually assumed that the fast reading times are a consequence of a shallow parse (due to, for example, good-enough processing, Ferreira, Bailey, & Ferraro, 2002) caused by cognitive overload. There can be good reasons to assume that shorter reading times are associated with increased difficulty, for example, when shorter reading times co-occur with lower comprehension accuracy (Van Dyke & McElree, 2006) or lower working memory capacity (Nicenboim et al., 2015; Nicenboim et al., 2016). However, the trade-off between reading times and comprehension accuracy is usually left underspecified.

There are two models that make explicit the relationship between reading times and retrieval accuracy, and even though they are sometimes not differentiated, they constitute two different implementations of the content-addressable cue-based retrieval mechanism. These models are the Lewis and Vasishth's (2005) *activation-based* model, and McElree's (2000) *direct-access* model. These models have different implications for retrieval processes in sentence comprehension. The activation-based model assumes a process that resembles a race model (Audley & Pike, 1965; Vickers, 1970), where evidence for each retrieval candidate is accumulated with different rates. This race determines both the latencies and the retrieval accuracy. By contrast, the direct-access model assumes that retrieval candidates have different levels of *availability*, which is the probability that a memory representation is retained. Availability determines only the accuracy of the retrieval and not the latency. In this model, a difference in latency between two conditions is a by-product of a mixture of directly accessed items, and retrievals that were initially incorrect, but they are reanalyzed leading to a correct retrieval.

The goal of this paper is to unpack the quantitative predictions of the activation-based and direct-access models by implementing them in a Bayesian hierarchical framework. This will allow us to compare their relative fit to a representative dataset and to assess their validity as models of retrieval that can account for similarity-based interference. We used a subset of the data from Nicenboim et al. (Submitted), which investigated similarity-based interference from the number feature using two relatively large-sample self-paced reading experiments. The data in this study include two dependent measures: (i) reading times for the critical region where retrieval from memory is assumed to occur, and (ii) accuracies in a comprehension task that targets specific dependency relations through a multiple choice task. This dataset is especially suitable for our modeling purposes because, apart from Van Dyke (2007), who also evaluated some of the dependencies, this is the only dataset that we are aware of that uses comprehension questions to directly assess the resolution of the dependencies. As explained in detail later, these two dependent measures (reading times and accuracy) are necessary for evaluating the models. We begin by describing the Nicenboim et al. study.

The Nicenboim et al. study

Nicenboim et al. (Submitted) used stimuli like (1). There were two conditions, high vs. low interference, which were assumed to affect the dependency between the subject (i.e., *Der Wohltäter* “The philanthropist”) and the verb (i.e., *begrüßt hatte* “had greeted”). In the high interference condition, two nouns intervened between these two co-dependents that had the same number marking as the target noun, the subject of the sentence, namely, singular marking. In the low interference case, the two intervening nouns had plural marking while the target noun remained singular. In German, the verb (i.e., *begrüßt hatte*) agrees in number with its subject; in the high interference condition, the retrieval cue set at the verb to seek out a singular noun would match three nouns. By contrast, in the low interference condition, only one noun matches this retrieval cue. Thus, reading time at the

critical region, the verb *begrüßt hatte*, provides an estimate of any interference effect.

Each target sentence was followed by a question that queried either the subject of the matrix verb (e.g., “sat”), the subject of the embedded verb (e.g., “had greeted”), or the object of the embedded verb. The possible answers were provided in multiple-choice format in pseudo-randomized order. For all the questions, participants had the option to answer “I don’t know”, when they did not remember or could not answer.

(1) a. HIGH INTERFERENCE

Der Wohltäter, der den Assistenten
The.sg.nom philanthropist, who.sg.nom the.sg.acc assistant (of)
 des Direktors begrüßt hatte, saß später im
 the.sg.gen director greeted had.sg, sat.sg later in the
 Spendenausschuss.
 donations committee.

‘The philanthropist, who had greeted the assistant of the director, sat later in the donations committee.’

b. LOW INTERFERENCE

Der Wohltäter, der die Assistenten
The.sg.nom philanthropist, who.sg.nom the.pl.acc assistants (of)
 der Direktoren begrüßt hatte, saß später im
 the.pl.gen directors greeted had.sg, sat.sg later in the
 Spendenausschuss.
 donations committee.

‘The philanthropist, who had greeted the assistants of the directors, sat later in the donations committee.’

Nicenboim et al. (Submitted) ran two experiments that showed an inhibitory effect of similarity-based interference for the retrieval of the subject (“the philanthropist”) when it shared the number feature singular with other competitor NPs (“the assistant”, “the director”), that is retrieval in high interference conditions took longer than in low interference conditions. Given that the auxiliary verb (*hatte* “had”) is morphologically marked as singular, the longer reading times at the auxiliary verb is consistent with

cue-based retrieval. Figure 1 shows the posterior distribution of the difference of reading times between conditions.




Figure 1. Posterior distribution of the estimated difference in reading times between conditions (in milliseconds) for the pooled data of Nicenboim, Engelmann, Suckow, and Vasishth (Submitted). The vertical line indicates the mean of the posteriors, the outer error bars demarcate the 95% credible intervals, and the inner error bars and filled section of the distributions the 80% credible interval.

Both the activation-based and the direct-access models make the correct predictions regarding the average behavior—both predict an inhibitory interference effect. However, these two accounts differ in the way that correct and incorrect retrievals occur, and these different underlying mechanisms can be investigated using the data from Nicenboim et al. (Submitted).

We first describe the models qualitatively, but in order to unpack the predictions of these two models we later provide a more formal presentation. We then evaluate the models quantitatively by examining the relationship between reading times and accuracy.

There are two main findings in the present study: First, the direct-access model provides a better fit to the data than the activation-based model. Second, we show that a variation of the activation-based model fits the data as well as the direct-access model, and also provides a reasonable model of the underlying generative process. A surprising aspect of the model selection is that it is driven by the difference in latencies for correct and incorrect retrievals, and not by the difference between interference conditions.

Overview of the activation-based and direct-access model

The activation-based model. The activation-based model as implemented by Lewis and Vasishth (2005) is a computational model of sentence processing in which dependencies of non-adjacent elements are created via a content-addressable cue-based retrieval mechanism. This model was realized in ACT-R (Anderson et al., 2004), which is a general cognitive architecture used to model a vast variety of cognitive phenomena. The activation-based model uses the declarative retrieval module of ACT-R, which has been shown to be able to account for many memory-related phenomena (e.g., Anderson, Bothell, Lebiere, & Matessa, 1998; Anderson & Reder, 1999; Van Rijn & Anderson, 2003). This means that sentence processing depends on the application of general cognitive principles to the specialized task of sentence parsing. Being a computational model, it provides quantitative predictions of retrieval speed and accuracy.

The predictions regarding interference, locality, and some antilocality effects of Lewis and Vasishth's (2005) original model have been investigated using simulations (Lewis & Vasishth, 2005; Vasishth et al., 2008). In addition, simplified versions of the model, which focused on certain aspects of the architecture and evaluated some of the assumptions of the original model, have also been used (e.g., Dillon, Mishler, Sloggett, & Phillips, 2013; Dillon et al., 2014; Kush & Phillips, 2014; Jäger et al., 2015; Vasishth & Lewis, 2006; Nicenboim et al., 2016; Engelmann, 2015; Parker & Phillips, 2016; Parker & Phillips, 2017; Engelmann, Jäger, & Vasishth, Submitted).

Crucially, the activation-based model provides an account of the relationship between reading times at the dependency resolution site and the accuracy of the resolution. This is so because dependency creation relies on the retrieval of the correct item from memory; in ACT-R terms, what is retrieved is a chunk. The chunk with the highest activation is retrieved and its activation level determines the retrieval time with higher activation leading to shorter times. We provide next an informal explanation of the key aspects of the activation-based model. We do this using example (1) from Nicenboim et al. (Submitted).

We show that the activation-based model can explain similarity-based interference effects, predicting inhibitory interference (i.e., an increase in processing difficulty) when a competitor NP matches the singular number feature of the target of retrieval.

The main assumptions of the model are that (i) words and phrases are encoded in memory as bundles of features (as in Nairne, 1990; Oberauer & Kliegl, 2006) that include lexical, semantic, and syntactic information, and that (ii) retrieval cues are used to identify the “correct” chunk from memory: If retrieval cues (which are feature specifications) match with the features of a chunk in memory, the chunk gets a boost in activation, and if cues mismatch a chunk’s features, activation is decreased. Such a mechanism would always retrieve the correct item. However, due to noise in the system, activation fluctuates randomly from trial to trial, so that despite a cue match with a target, a competitor could have higher activation and could end up being retrieved. An alternative possibility is that all candidate chunks in memory fall below a retrieval activation threshold (a parameter in ACT-R); in this case, retrieval would fail.

As an example, consider the auxiliary verb (*hatte*, “had.sg”) of (1). This is the region where an interference effect was seen in Nicenboim et al. (Submitted). In both sentences (1a) and (1b), there is a dependency between this verb and its subject (“the philanthropist”), and the only difference between the sentences is that the intervening NPs (“the assistant/s”, “the director/s”) appear in singular in (1a) and in plural in (1b). The activation-based model assumes that the feature information of each item such as category, case, number, gender, and so forth is encoded in memory. When the embedded verb (*begrüßt hatte*, “greeted had.sg”) is being read, an attempt is made to retrieve the subject. The verb provides cues such as *NP*, *nominative* (notice that case is encoded in the determiner of the NP in German), *singular*, among others features required from the target of the retrieval. For each cue, a limited amount of activation (called the maximum associated strength or *MAS*) is spread among the target and the competitors that are stored in memory. The *MAS* determines the strengths of association from each cue to each

item in memory. This strength of association represents how uniquely the cue identifies a target. This means that in the low interference condition (1b), the strength of association of the singular cue with the target is determined by the maximum activation associated with this cue (since the cue fully identifies the item). In the high interference condition (1a), however, the target (“the philanthropist”) and the competitors (“the assistant”, “the director”) will be assigned some smaller part of the maximum activation (this is the so-called fan effect, for details, see Anderson & Reder, 1999), and thus their strength of association of *singular* will be smaller than the maximum activation. This is regardless of the fact that in both conditions, there is another cue that uniquely identifies the target, namely, being nominative: In both conditions the target also receives activation due to the strength of association with the cue nominative. This means that (i) the target in high interference conditions would receive (on average) less total spreading activation than the target in low interference conditions, (ii) the competitors in high interference conditions would receive (on average) more total spreading activation than the competitors in low interference conditions, and (iii) the target would receive (on average) more total spreading activation than the competitors. See Figure 2 for a schematic that illustrates this. Given that higher activation leads to shorter retrieval times, the way activation is shared among the NPs leads to inhibitory interference, that is, longer retrieval times and lower accuracy in high interference conditions in comparison with low interference ones. Another less studied consequence is that the incorrect retrieval of the competitors should take longer than the correct retrieval of the target. We will return to this last point in the discussion section.

Given that both latency and probability of successful retrieval are affected by activation, we will show later that the retrieval process is similar to a race of accumulators (among many others: Audley & Pike, 1965; Vickers, 1970; Usher & McClelland, 2001): Each item in memory is assigned an *accumulator of evidence* for its retrieval, where the activation of each item acts as the rate of accumulation. The accumulator that reaches the threshold of evidence first determines which item is retrieved and with which latency. This

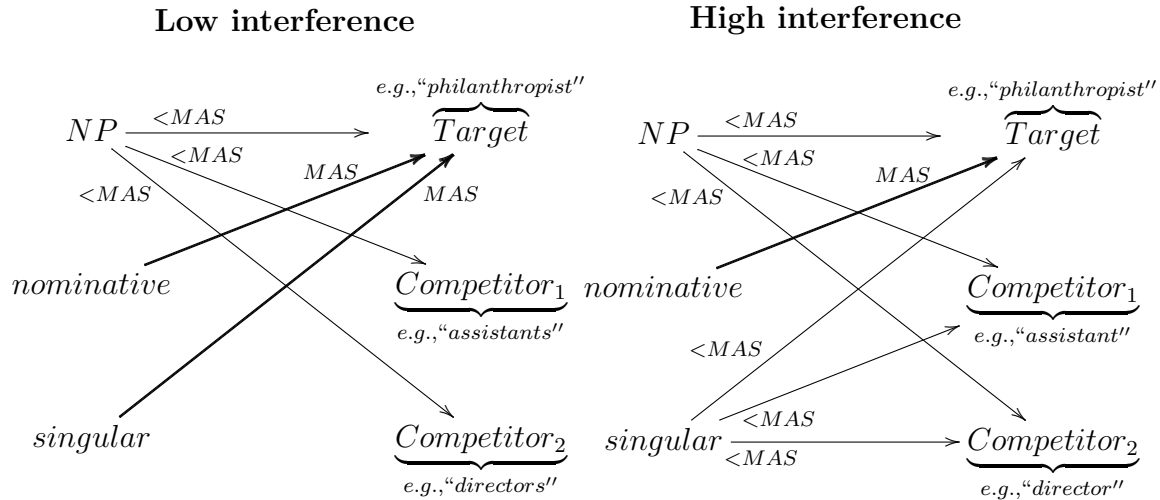


Figure 2. Graph showing the associations between the cues NP, nominative, and singular, and the target and two competitor NPs. The width of the arrow represents the strength of association, but this strength is weighted, so two identical associations may assign different amount of activation depending on the cue to which they belong.

places the model under a sequential sampling framework (such as the drift diffusion model: Ratcliff, 1978; Ratcliff & McKoon, 2008; the leaky competitive accumulator: Usher & McClelland, 2001; linear deterministic models: Brown & Heathcote, 2008, among others; Heathcote & Love, 2012). Furthermore, some of the assumptions of ACT-R can allow us to frame the retrieval process as one of simplest accumulator models: the lognormal race model with a single variance for the noise associated with target, competitors, and failure accumulators (Heathcote & Love, 2012; Rouder, Province, Morey, Gomez, & Heathcote, 2014). A variant that we will consider later is a model with two separate variances, one for the target accumulator, and one for the competitors and failure accumulators.

The direct-access model. It should be noted that a content-addressable system does not necessarily entail a race between items in memory, and there are other models that are also compatible with a content-addressable cue-based retrieval mechanism. The cue-based retrieval model proposed in Van Dyke and McElree (2006) is based on McElree and colleagues' previous work (e.g., McElree, 2000; McElree, Foraker, & Dyer, 2003) and, while it does not assume a race model, it shares with the activation-based model some of the assumptions of cue-based retrieval: Words and phrases are also encoded in memory as

feature bundles, and retrieval cues are used to distinguish the target from the competitors. Whereas here the cues are combined multiplicatively (as proposed in the Search of Associative Memory (SAM) model of Raaijmakers & Shiffrin, 1980; Gillund & Shiffrin, 1984),¹ and in the activation-based model the cues are combined additively, the result is similar: The probability of retrieving a particular item from memory given the retrieval cues is a function of the degree of the match between the cues and the item, reduced by the degree to which the cues match other competitor items in memory. However, in contrast with the activation-based account, cues are supposed to enable *direct access* to relevant memory representations. This means not only that there is no serial search between items in memory, but that the distribution of access time is independent of the degree of match between item and cue, and regardless of the quality or strength of the representation of the item in memory (McElree, 2000). As a consequence of the direct access, in this model, the probabilities of retrieval are affected by inhibitory interference, resulting in a pattern similar to the one described for the activation-based model: (i) the probability of retrieving the target in high interference conditions would be lower than in low interference conditions, (ii) the probability of retrieving one of the competitors in high interference conditions would be higher than in low interference ones, and (iii) the probability of retrieving the target would be higher than the probability of retrieving one of the competitors.

It is not uncommon, however, that both poorer accuracy and longer reading times associated with similarity-based interference are taken as evidence for McElree's (2000) direct-access model as well as for Lewis and Vasishth's (2005) activation-based model. For example, Van Dyke and McElree (2006) write:

The current experiment supports a retrieval-based account of interference effects in sentence processing, one that is compatible with the hypothesis that a

¹Van Dyke and McElree (2011) are explicit, however, in that it could also be that cues are combined linearly, but with weights that are different enough so that certain cues, such as syntactic cues, have a more prominent role.

cue-based retrieval mechanism mediates the creation of grammatical dependencies during parsing. One such mechanism has been proposed in Van Dyke and Lewis (2003; see also Lewis and Vasishth, 2005; Van Dyke, 2002), in which parsing success depends on the extent to which required constituents can be retrieved from working memory. On this account grammatical heads provide retrieval cues that are used to access previously stored items via a content-addressable retrieval process (McElree, 2000, 2006; McElree, Foraker, & Dyer, 2003).

A slowdown in self-paced reading or eyetracking-while-reading can also be taken as evidence for direct access, since processing speed may be affected by differences in the likelihood of recovering an item from memory (McElree, 1993; McElree et al., 2003). This is because McElree (1993) assumes that after a misretrieval, that is, an incorrect or failed retrieval, the parser can often backtrack to reprocess the retrieval and reach the appropriate analysis. This would mean that a correct interpretation of a dependency could be arrived at because the correct dependent was retrieved at the first attempt or, alternatively, because a wrong dependent was retrieved initially but the parser backtracked and retrieved the correct one. Given that backtracking should take some additional time, latencies associated with the correct responses would be a mixture of fast directly accessed dependents and retrievals slowed down due to the time needed for backtracking. Since interference adversely affects retrieval probabilities, the proportion of errors would be higher in high interference conditions. This would entail a higher proportion of backtracking and hence slower latencies in the mixture of correct responses and on average, high interference conditions would show longer reading times than low interference conditions.

Since the assumed constant distribution of retrieval times may not be observable in reading for comprehension, evidence compatible with the direct-access model but incompatible with the activation-based model comes only from findings of speed-accuracy

trade-off (SAT) procedure on rapid grammaticality judgment task (e.g. McElree, 2000; McElree et al., 2003; Van Dyke & McElree, 2011). In this task, participants need to judge a sentence as either grammatical or ungrammatical, and their judgment process is interrupted with a cue to respond (typically a tone) after varying amounts of time (Reed, 1973; Wickelgren, 1977; and see also: Foraker & McElree, 2011). However, this is a meta-linguistic task and it would be desirable if independent support for the direct-access model could be found with a reading-for-comprehension task. In addition, the conclusion that there is a constant distribution of retrieval times requires arguing for a null result in one of the parameters of the SAT model. For SAT procedures, accuracy is modeled as a function of three parameters corresponding to the three phases of the SAT curve: (i) the asymptotic level of performance, (ii) the intercept or the point in time where performance is different from chance, and (iii) the rate at which accuracy grows from chance to asymptote. While the presence of the effect in the asymptote, such that an increase on interference lowers the asymptote, is evidence for the reduction of the probability of accessing the target, the lack of evidence for changes in the rate and intercept must be taken as evidence for no effect on the speed of the retrieval. It could be, however, that the differences were too small to be detected.

Since both models give virtually identical predictions for non-SAT (self-paced reading or eye-tracking) experiments when comparing only reading time averages across conditions, a slowdown at high interference conditions has been taken as evidence for both models. However, the two models assume different relationships between retrieval times and responses. It is important to assess the fit to the data of each model, since each one is compatible with different memory retrieval mechanisms.

Implementation of the models

In order to distinguish between the activation-based and direct-access models, we implemented them as hierarchical Bayesian models. Implementing these models affords

several advantages: (i) we can investigate (through posterior predictive checking, explained below) whether the data could have been generated by the models; (ii) we can determine how each model’s parameters were affected by interference effects; and (iii) the quality of fit of the two models can be compared using cross-validation. It is important to note that we are using Bayesian methods as a flexible and interpretable way of extending models of cognitive processes (Lee, 2011; Shiffrin, Lee, Kim, & Wagenmakers, 2008), and this approach is orthogonal to the question of whether the mind does or does not do Bayesian inference (see Feldman, 2016, for a critical review on how probabilistic models of perception and cognition should work).

The benefits of using hierarchical Bayesian modeling are two-fold: (i) The models incorporate the general advantages of Bayesian inference, such as the use of credible intervals instead of confidence intervals, and the possibility of fitting complex non-linear models (see Nicenboim & Vasishth, 2016, for an extended discussion), and (ii) hierarchical models allow us to take both between- and within-group variance into account and pool information via shrinkage (Gelman, Hill, & Yajima, 2012). This means that we avoid overfitting the data and at the same time we avoid aggregating data and losing valuable information about group-level variability (Gelman & Hill, 2007). In the next section we provide a more formal presentation of the assumptions and details of the implementation of both models than the one given in the introductory overview.

The activation-based model as a lognormal race

In order to assess the fit of the activation-based account to the experimental data, we implemented it as a lognormal race of accumulators with a shift parameter and a single variance for the noise of all accumulators (Rouder et al., 2014), as explained further below. The retrieval in ACT-R can be thought of as a decision processes, where target and competitors stored in memory accumulate evidence until the first chunk reaches a certain value and is retrieved. Activation can be linked to evidence by assuming that it represents

the rate of its accumulation (in a similar way as assumed by van Maanen, van Rijn, & Taatgen, 2011). This is so because activation in ACT-R represents the probability of the retrieval, and it is boosted by processes that increase evidence such as matching cues, previous retrievals and so forth, while it is penalized by processes that decrease the evidence such as mismatching features and decay.

Since the chunk with the highest activation is retrieved, the equation that determines the latency of the retrieval in ACT-R (1) ignores the activation of the other chunks. Our implementation, however, assumes that there is a *potential* retrieval time, or *finishing time* in the race, t_c for each candidate c in memory that depends on its activation, A_c . This is the time it would have taken for the chunk to be retrieved given its activation, as shown in Equation (2).

$$Latency \propto e^{-\arg \max_c (A_c)} \quad (1)$$

$$t_c \propto e^{-A_c} \quad (2)$$

The race model is implemented in the following way: Since the noise component in A_c is assumed to be normally distributed (Lebiere, Anderson, & Reder, 1994) and affects the activations of all the chunks to the same extent,² for each trial, the finishing time of each chunk, t_c , is sampled from a lognormal distribution with the same standard deviation σ , and the fastest chunk in a given trial would be the one retrieved (i.e., the chunk with the lowest t_c in a given trial). We account other aspects of processing (e.g., lexical access) with a parameter γ .

²The noise component is sometimes approximated to have a logistic distribution for convenience, see, for example Lebiere (1999).

$$t_c \sim e^{normal(-\mu_c + \gamma, \sigma)} \Rightarrow \log(t_c) \sim normal(-\mu_c + \gamma, \sigma) \quad (3)$$

$$\Leftrightarrow t_c \sim lognormal(-\mu_c + \gamma, \sigma) \quad (4)$$

where

$$A_c = \mu_c + \epsilon \quad (5)$$

$$\epsilon \sim normal(0, \sigma) \quad (6)$$

The only observable data for every trial are (i) the answer of the comprehension question at the multiple choice task, w , (we assume that when the question asks about the subject of the embedded verb, the answers correspond to the chunk retrieved from memory, i.e., the winner of the race, modulo offline distractions), and (ii) the reading times at the site of the retrieval (the auxiliary verb “had.sg”). The reading times will include the retrieval time, $t_{c=w}$, of the “winner” chunk, and the time taken for other processes. Given the evidence that distributions of reaction times are shifted (Rouder, 2005; Nicenboim et al., 2016), we assume a lower bound, ψ , which represents changes in peripheral aspects of processing, such as encoding or motor execution (Rouder, 2005).

$$RT_{c=w} \sim \psi + lognormal(-\mu_c + \gamma, \sigma) \quad (7)$$

Since the observed reading times, as shown in Equation (7), are associated with the maximum activation (on a specific trial), the observable data constrain the unobserved finishing times of all the other non-selected choices c which are not the winner; these finishing times must be slower than the finishing time of the winner w . We can write this

as Equation (8).

$$t_{\forall c, c \neq w} > t_{c=w} \quad (8)$$

Since we are not interested in the specific value of μ_c , or γ , but in learning from the model (i) whether the retrieval process resembles a race of accumulators, and (ii) the effect of number interference on the target and competitors, we fit the reading times, RT , as a function of α_c and an arbitrary constant, b , such that $b - \alpha_c = -\mu_c + p$. By setting b large enough (to 10, for example), we ensure that α_c is strictly positive for ease of interpretation: a higher positive number corresponds with a higher rate of accumulation. Rouder et al. (2014) show that without further assumptions, thresholds and accumulation rates cannot be disentangled in the lognormal race model. To estimate the thresholds in a lognormal race model, Heathcote and Love (2012) assumed that both the rate of accumulation ν and the thresholds η were lognormally distributed, so that the finishing times, y , were distributed in the following way $y \sim \text{lognormal}(\mu = \mu_\eta - \mu_\nu, \sigma = \sqrt{\sigma_\eta^2 + \sigma_\nu^2})$. If the thresholds are fixed at some arbitrary point b , then $\sigma = \sigma_\nu$ (since $\sigma_\eta = 0$) and the rate of accumulation $\mu_\nu = b - \mu$, thus we can interpret α_c as the rate of accumulation associated with each chunk, and we can rewrite Equation (7) in the following way:

$$RT_{c=w} \sim \psi + \text{lognormal}(b - \alpha_{c=w}, \sigma) \quad (9)$$

On every trial, l , we can estimate $\alpha_{l,c=w}$ from the observed RT , and we can constrain the possible values of $\alpha_{l,\forall c, c \neq w}$ from the values of $t_{l,\forall c, c \neq w}$ that could not be possible on a given trial.

Given that shifts vary across participants but tend not to vary with experimental manipulation (Rouder, 2005), we assume a certain shift for every participant i , while $\alpha_{l,c}$, the activation together with nuisance parameters, will vary by participant i and by

experimental item j .

$$RT_{l,i,j} \sim \psi_i + \text{lognormal}(b - \alpha_{l,i,j,c=w}, \sigma) \quad (10)$$

In standard ACT-R, if the activation is below a certain threshold, T , the retrieval fails with a latency that is proportional to e^{-T} (Lebiere et al., 1994). To avoid a deterministic latency, we assign an accumulator to the possibility of failure, which acts as a noisy timer, and its timeout depends on its parameter $\alpha_{c=failure}$. In this way, the retrieval threshold can be also thought as a chunk that competes for activation (Van Rijn & Anderson, 2003).

Figure 3 summarizes the parameters and the process for two chunks. The lognormal race model is sometimes called a *ballistic* or *deterministic* race model because there is no within-choice noise (as in, for example, the drift diffusion model): once a rate of accumulation is set for a given accumulator, it will determine its time to the threshold; this is represented in the lower part of Figure 3 by the straight lines. This, however, does not make the process deterministic, since the rate of accumulation changes from trial to trial.

In the implementation of the activation-based model, the hierarchical structure of this model is embedded in each parameter α_c associated with each chunk c (including one representing the failure) that is allowed to vary by condition (high/low interference) and includes by-participants and by-experimental-items intercepts and slopes (which are also allowed to be correlated). This means that the model can account, for example, for an NP of a certain experimental item being more semantically plausible as a retrieval candidate than other NPs, by simply adjusting the by-item intercept of the accumulators associated with each NP. See Appendix A for the details of the Bayesian model.

The direct-access model as a mixture model

Next, we present an implementation of a direct-access model. This is a Bayesian hierarchical implementation of the cue-based retrieval model proposed by Van Dyke and McElree (2006) based on McElree and colleagues' previous work (e.g., McElree, 2000;

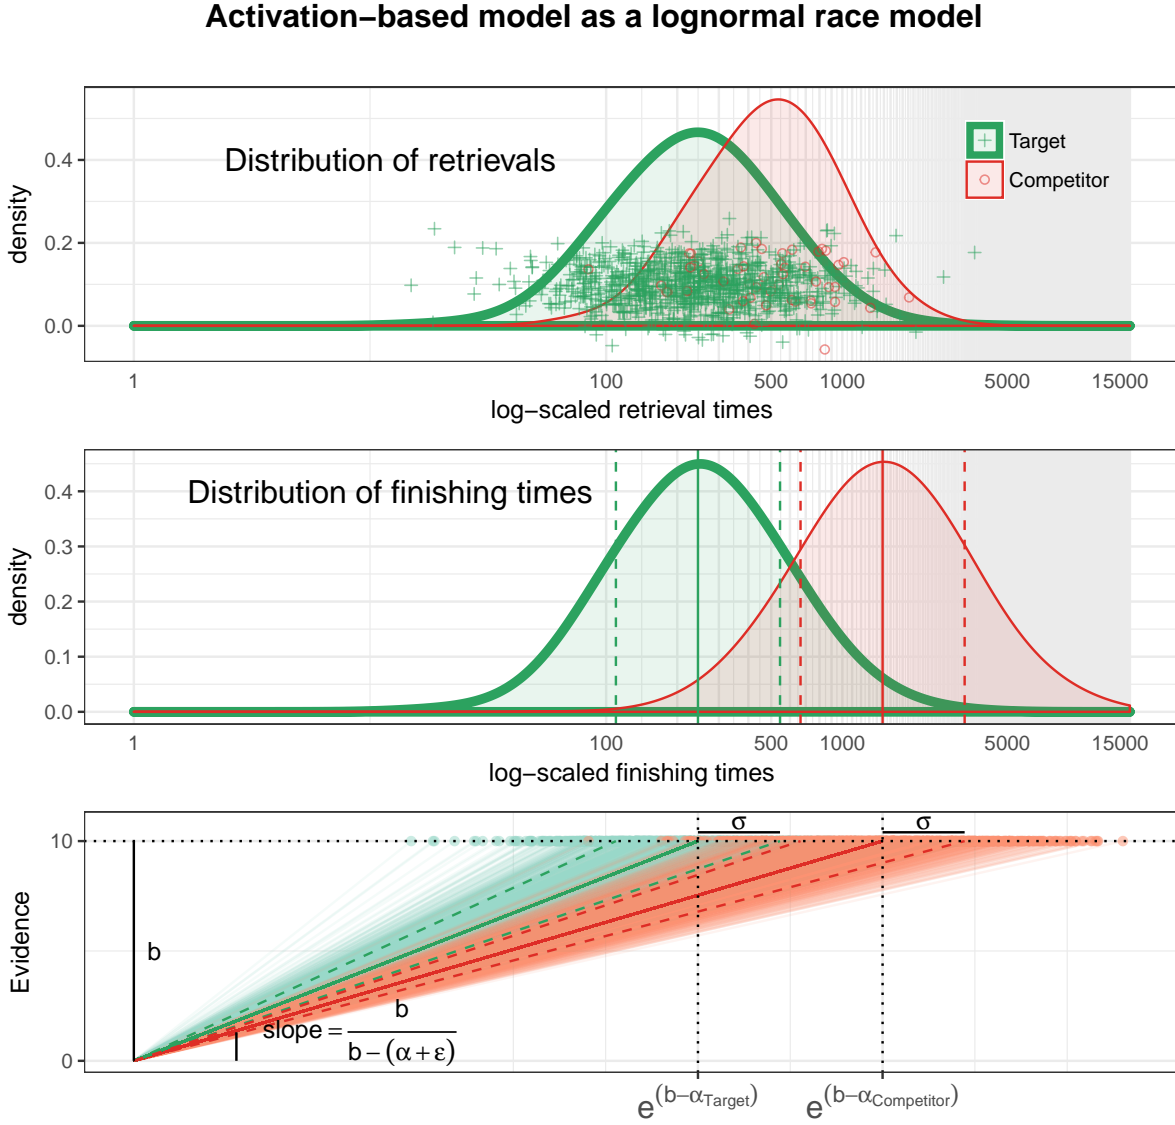


Figure 3. The figure depicts how the distribution of retrievals is generated from the activation-based model as a lognormal race model. The bottom figure depicts the parameters of the activation-based model, the full lines in green and red are the mean finishing times $t_{c=Target}$ and $t_{c=Competitor}$, and the broken lines are finishing times one standard deviation away from the mean. The middle figure shows the distributions of finishing times for target and candidate; since every chunk is associated a potential finishing time, t , both distributions have the same number of elements. The top-most figure shows the distribution of retrieval times (adding the shift parameter ψ would transform it to reading times); since only the winning chunks are retrieved, the distribution of retrieval times for targets has more elements than the one of the competitors. Notice that the retrieval times are faster than the finishing times: this is so because when a chunk has a very long finishing time in a given trial, it is very likely that its competitor will be faster.

McElree et al., 2003). Since the original model has never been implemented computationally, it is underspecified in some respects. We therefore made some assumptions in the model to spell these details out; these are described below.

In order to account for differences in reading times, the direct-access model assumes that, in some proportion of the cases, the parser is able to backtrack a misretrieval and to access the target candidate, taking some extra time. Thus the reading times associated with the correct retrievals are a mixture of directly accessed as well as backtracked retrievals as shown in Figure 4. This means that the probability of the final retrieval of a certain chunk, P_r (which should be equivalent to the proportion of responses given at the multiple choice task modulo offline distractions), is affected by the probability of backtracking, θ_b , together with the probability of either an initial retrieval of the target, θ_{Target} , an initial retrieval of one of the competitors, $\theta_{Competitor_{1,2}}$, or an initial failure, $\theta_{failure}$, in the following way:

$$P_r(Target) = \theta_{Target} + (1 - \theta_{Target}) \cdot \theta_b \quad (11)$$

$$P_r(Competitor_{\{1,2\}}) = \theta_{Competitor_{\{1,2\}}} \cdot (1 - \theta_b) \quad (12)$$

$$P_{failure} = \theta_{failure} \cdot (1 - \theta_b) \quad (13)$$

where

$$\theta_{Target} + \theta_{Competitor_1} + \theta_{Competitor_2} + \theta_{failure} = 1 \quad (14)$$

The core assumption of the direct-access model is that retrieval takes the same time on average, t_{da} , regardless of the availability of the to-be-retrieved chunk. This is in contrast with the activation-based model, but also in contrast with the SAM framework of (Gillund & Shiffrin, 1984), where the retrieval time depends on the match between cues

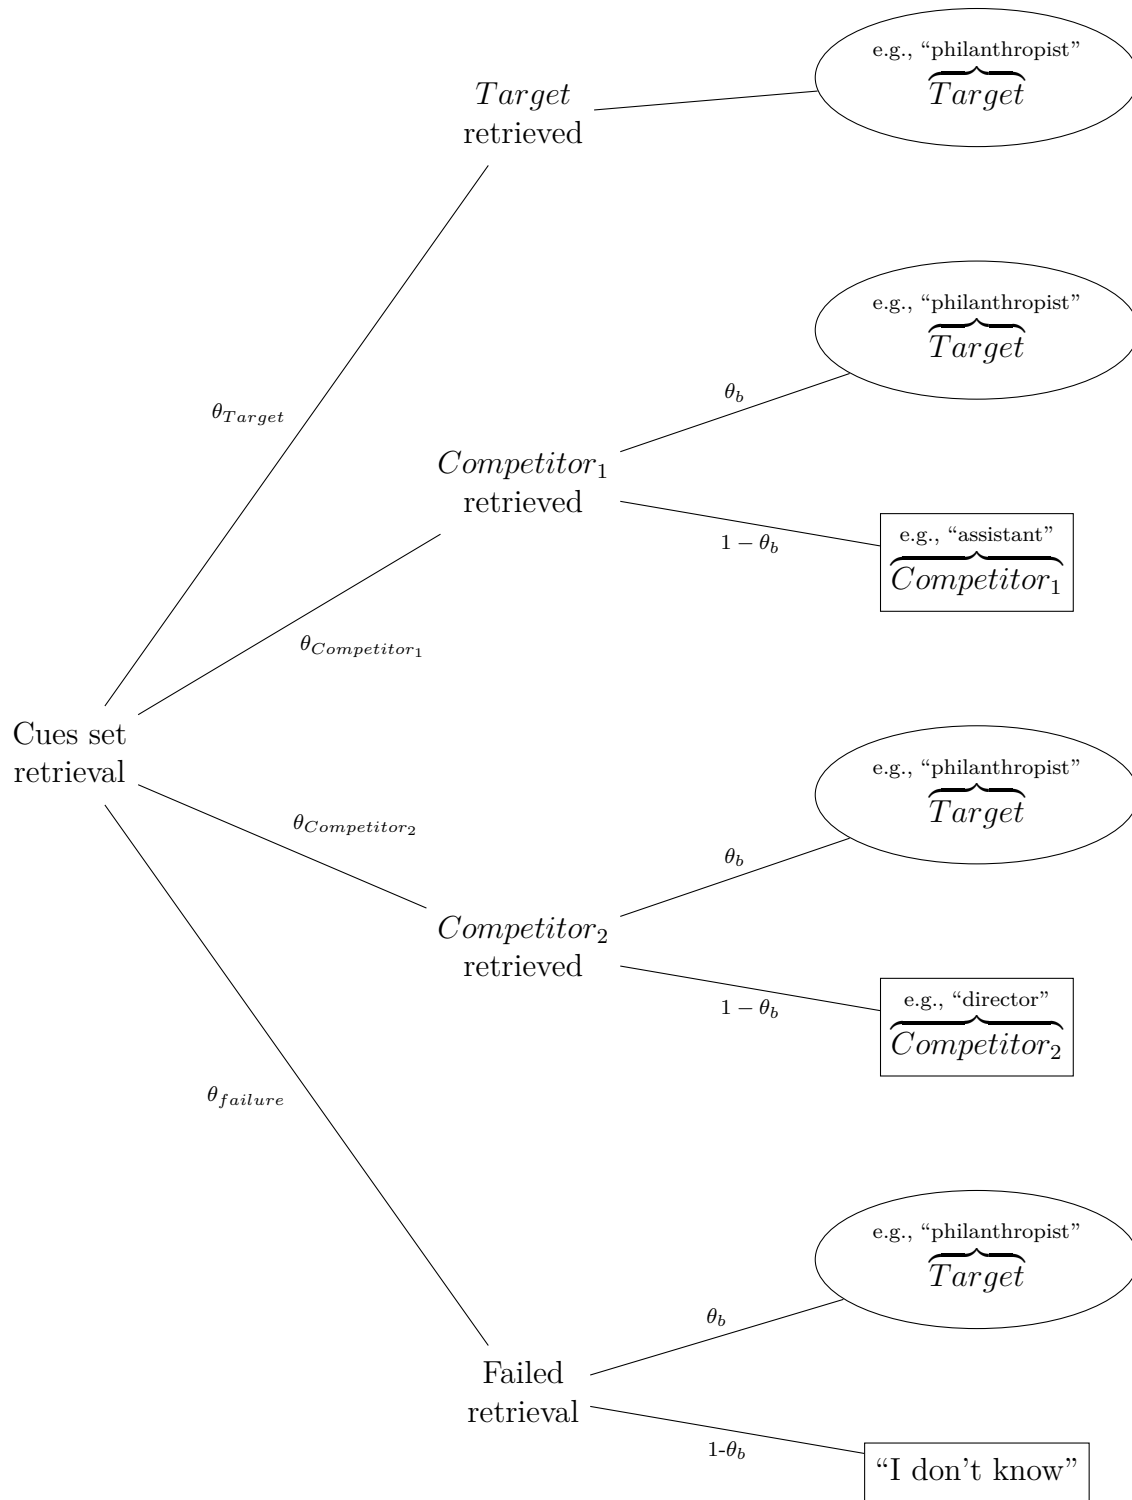


Figure 4. Graph showing the different responses as a categorical distribution with correct responses inflated by the probability of backtracking, θ_b . Correct responses are inside ellipsis and incorrect ones inside rectangles.

and features of the chunk. The implications for the direct-access model is that there will be two different distributions of reading times: a distribution associated with the incorrect responses and a distribution associated with the correct ones. An incorrect response is given by a participant, only when the wrong chunk is retrieved and there is no backtracking and repair process; see Figure 4. In this case, the reading times at the retrieval site will only include the time needed for the direct access, t_{da} , together with the time taken for other processes, γ , and normally distributed noise with standard deviation σ ; crucially, this time is independent of the level of interference. We assume, as before, that the noise and the time taken for other processes are added to the location of the lognormal distribution; in addition, reading times are assumed to be shifted by some minimum amount ψ of time that represents the lower bound of the process. Thus, we can assume that the reading times at the retrieval site for each trial l that are associated with an incorrect response (i.e., the retrieval of a competitor or a failed retrieval) have a shifted lognormal distribution, where the t_{da} depends on each participant i , and experimental item j , but not on the experimental condition or the identity of the chunk retrieved.

$$RT_{incorrect,l,i,j} \sim \psi_i + \text{lognormal}(t_{da,i,j} + \gamma_{i,j}, \sigma) \quad (15)$$

For correct responses, the reading times depend on whether there is a repair process or not; see Figure 4. This entails that the distribution of reading times is a mixture of two components: The first one is associated with chunks correctly retrieved at the first attempt as shown in the first line of Equation (16), and it is identical to the distribution of incorrect responses. The second one is associated with incorrect responses that are backtracked and repaired, and it includes the direct-access times, t_{da} , together with the time it takes to backtrack and do a reanalysis, t_b , as shown in the second line of Equation (16). (We derive

the exact proportions in Appendix B.)

$$RT_{correct,l,i,j} \sim \psi_i + \begin{cases} \text{lognormal}(t_{da,i,j} + \gamma_{i,j}, \sigma) & , \text{ if the first try is correct} \\ \text{lognormal}(t_{da,i,j} + t_{b,i,j} + \gamma_{i,j}, \sigma) & , \text{ otherwise} \end{cases} \quad (16)$$

As with the activation-based model, we are not interested in γ , and we define $T_{da,i,j} = t_{da,i,j} + \gamma_{i,j}$. In the implementation of the model, we estimate the probability of each retrieval and the effect of interference on the retrieval probability using a multi-logit regression (or categorical distribution with the parameters on the logit scale). This is achieved by assigning a hierarchical structure to the parameters of the multi-logit regression which vary by condition (high/low interference) and include by-participants and by-experimental-items intercepts and slopes which have one correlation matrix for participants and one matrix for experimental items. Furthermore, we assign a hierarchical structure also to T_{da} and t_b , which are composed by by-participant and by-experimental-item varying intercepts. To allow for correlations between the direct access time and the backtracking time, we included also one correlation matrix for participants and one matrix for experimental items. This means that we assume that latencies should not be affected by retrieval probabilities. See Appendix B for more details.

Evaluation of the activation-based and direct-access models

Application to data from a self-paced reading experiment

We fit the models to a subset of the data from Nicenboim et al. (Submitted). This work reports two self-paced reading studies: an exploratory and a confirmatory study with a total of 183 participants investigating similarity-based interference. Both studies included the same 60 experimental items, where a high interference condition with two competitor nouns sharing the singular feature with the target of the retrieval (e.g., “the director”) is contrasted with a low interference condition with two plural competitors, see (1), repeated

here for convenience as (2). In order to encourage attentive reading, after every sentence participants answered a multiple choice questions querying different aspects of the sentence: for example, (3).

(2) a. HIGH INTERFERENCE

Der Wohltäter, der den Assistenten
The.sg.nom philanthropist, who.sg.nom the.sg.acc assistant (of)
 des Direktors **begrüßt hatte**, saß später im
 the.sg.gen director **greeted had.sg**, sat.sg later in the
 Spendenausschuss.
 donations committee.

‘The philanthropist, who had greeted the assistant of the director, sat later in the donations committee.’

b. LOW INTERFERENCE

Der Wohltäter, der die Assistenten
The.sg.nom philanthropist, who.sg.nom the.pl.acc assistants (of)
 der Direktoren **begrüßt hatte**, saß später im
 the.pl.gen directors **greeted had.sg**, sat.sg later in the
 Spendenausschuss.
 donations committee.

‘The philanthropist, who had greeted the assistants of the directors, sat later in the donations committee.’

(3) Wer hatte jemanden begrüßt?

Who had greeted someone?

- | | |
|--------------------------------|------------------------|
| a. der/die Wohltäter (CORRECT) | c. der/die Direktor/en |
| the philanthropist(s) | the director(s) |
| b. der/die Assistent/en | d. Ich weiß es nicht |
| the assistant(s) | I don’t know |

For the current study, we pooled the data of both self-paced reading experiments, keeping only the sentences with questions that queried the subject of the embedded verb,

as in (3). This was done because cue-based retrieval models predict that interference will affect only retrievals where the cue is relevant. This left us with 20 sentences for each participant. For each sentence we used two dependent variables: (i) the time taken for reading the auxiliary verb *hatte* (“had.sg”), where Nicenboim et al. (Submitted) found evidence for interference effects in reading times, and (ii) the response given at the multiple choice task.

In order to fit the models, however, we need *retrieval times* and the *retrieved argument*. Thus we assume that all things being equal (as in an experimental setting), (i) reading times at the auxiliary verb include the retrieval time (as well as other theoretically uninteresting processes), and (ii) participants’ answers are informative of the dependency completion they carried out.

We fit the models using *rstan* package (Stan Development Team, 2017, version 2.15.1) in R (R Core Team, 2015) with four chains and 3000 iterations, half of which were the burn-in or warm-up phase. In order to assess convergence, we verified that the potential scale reduction factors, \hat{R} s, of the parameters were close to one, and we also visually inspected the chains (Gelman, Carlin, et al., 2014, section 11.4). This indicates that the chains have mixed and they are transversing the same distribution. When needed, we also increased the maximum tree-depth and the adaptation parameter δ of the sampler to eliminate divergent transition and achieve convergence. We also verified that we could recover the parameters from the models using fake-data simulation (Gelman & Hill, 2007), see Appendix C.³

Posterior predictive checking

We use posterior predictive checking to examine the descriptive adequacy of the models (Shiffrin et al., 2008; Gelman, Carlin, et al., 2014, Chapter 6), that is, the observed data should look plausible under the posterior predictive distribution. The posterior

³Data and code can be downloaded from: www.ling.uni-potsdam.de/~nicenboim/code/code-data-retrieval-models.zip

predictive distribution is composed of 6000 datasets (one for each iteration) that the model generates based on the posterior distributions of its parameters. In other words, given the posterior of the parameters of the model (which are based on the current data), the posterior predictive distribution shows how other data may look like. Achieving descriptive adequacy means that the current data could have been predicted with the model. While passing a test of descriptive adequacy is not strong evidence in favor of a model, a major failure in descriptive adequacy can be interpreted as strong evidence against a model (Shiffrin et al., 2008). Thus, posterior predictive checking is an important sanity check to assess whether the model behavior is reasonable (see Gelman, Carlin, et al., 2014, for further discussion)

Given that the main difference between the activation-based model and the direct-access model is in the way they account for the relationship between retrieval probability and latencies, for each of the 6000 datasets generated by the models, we calculate the means and .1–.9 quantiles of the reading times associated with each response, as well as the mean proportion of responses given. We represent this graphically using violin plots (Hintze & Nelson, 1998): the width of the violin plots represents the density of the predicted means (or quantiles). The observed mean (or quantile) of the data is represented with a cross. If the data could plausibly have been generated by the model, we would expect the crosses to be inside the violin plots.

Estimation of relevant parameters

In addition to fitting the data, the models include parameters that can be interpreted and can give support (or falsify) some assumptions of the effect of interference under the two presented models. In particular, the following inequalities represent the expected relationships between activations for the activation-based model and how they are affected by interference (with $|HI$ indicating under high interference, and $|LI$ under low interference):

$$\alpha_{Target|HI} < \alpha_{Target|LI} \quad (17)$$

$$\alpha_{Target} > \{\alpha_{Competitor_1}, \alpha_{Competitor_2}, \alpha_{Failure}\} \quad (18)$$

$$\alpha_{Competitor_{\{1,2\}}|HI} > \theta_{Competitor_{\{1,2\}}|LI} \quad (19)$$

and similarly for the probability of retrieval at the first attempt in the direct-access model:

$$\theta_{Target|HI} < \theta_{Target|LI} \quad (20)$$

$$\theta_{Target} > \{\theta_{Competitor_1}, \theta_{Competitor_2}, \theta_{Failure}\} \quad (21)$$

$$\theta_{Competitor_{\{1,2\}}|HI} > \theta_{Competitor_{\{1,2\}}|LI} \quad (22)$$

In addition, the direct-access model includes the probability of backtracking, θ_b , and the time needed for backtracking, t_b .

We provide the estimates of the previous key parameters (or relations between key parameters) with their credible intervals.

Cross-validation

We also compared the models using cross-validation, since the descriptive adequacy can also be achieved by a model that is too flexible and can generate too many different results. The idea behind cross-validation is to assess the accuracy the model would have in making predictions for new data, that is the expected predictive performance.

We compare the expected predictive performance of the models (Gelman, Hwang, & Vehtari, 2014) using k-fold cross-validation (Vehtari & Ojanen, 2012; Vehtari, Gelman, & Gabry, 2017), with k set to ten. We calculated the k-fold cross-validation by first splitting the data into ten subsets (or folds) and then using each subset as the validation set, while

the remaining data were used for parameter estimation. We partitioned the data into subsets by pseudo-randomly permuting the observations, and then systemically dividing them into 10 subgroups; we ensured that each group contained similar number of observations for each subject (this was meant to avoid the situation where most of the data of a certain subject is left out due to chance). The estimate of the expected log pointwise predictive density (\widehat{elpd}) for a new dataset (i.e., the sum of the expected log pointwise predictive density of each observation) can be used as a measure of predictive accuracy for the total number of observation (N) taken N/k at a time; \widehat{elpd} can be transformed to deviance scale by multiplying it by minus two, providing a fully Bayesian alternative to AIC (Akaike Information Criterion; Akaike, 1974) or DIC (Deviance Information Criterion; Spiegelhalter, Best, Carlin, & Van Der Linde, 2002).

Results

Activation-based model.

Posterior predictive check. The posterior predictive check reveals that the model is inadequate for predicting some key characteristics of the data. Figure 5(a) shows that the model predicts shorter times for reading the auxiliary verb when the correct response is given and longer times for reading the auxiliary verb when an incorrect answer is given. In other words, the model underestimates the retrieval time of the correct dependent and overestimates the retrieval time of the competitor NPs, or the timeout. Figure 5(b) also shows a slight misfit for the predicted accuracy: the model tends to underestimate the proportion of correct responses and to slightly overestimate the proportion of incorrect ones. Furthermore, Figure 5(c) reveals that the fit is especially poor for the second half of the quantiles.

Estimation of relevant parameters. The key parameters and relationships between parameters of the activation-based model are summarized in Figures 6(a) and (b). Figure 6(a) shows caterpillar plots of the posterior distributions for the rates of

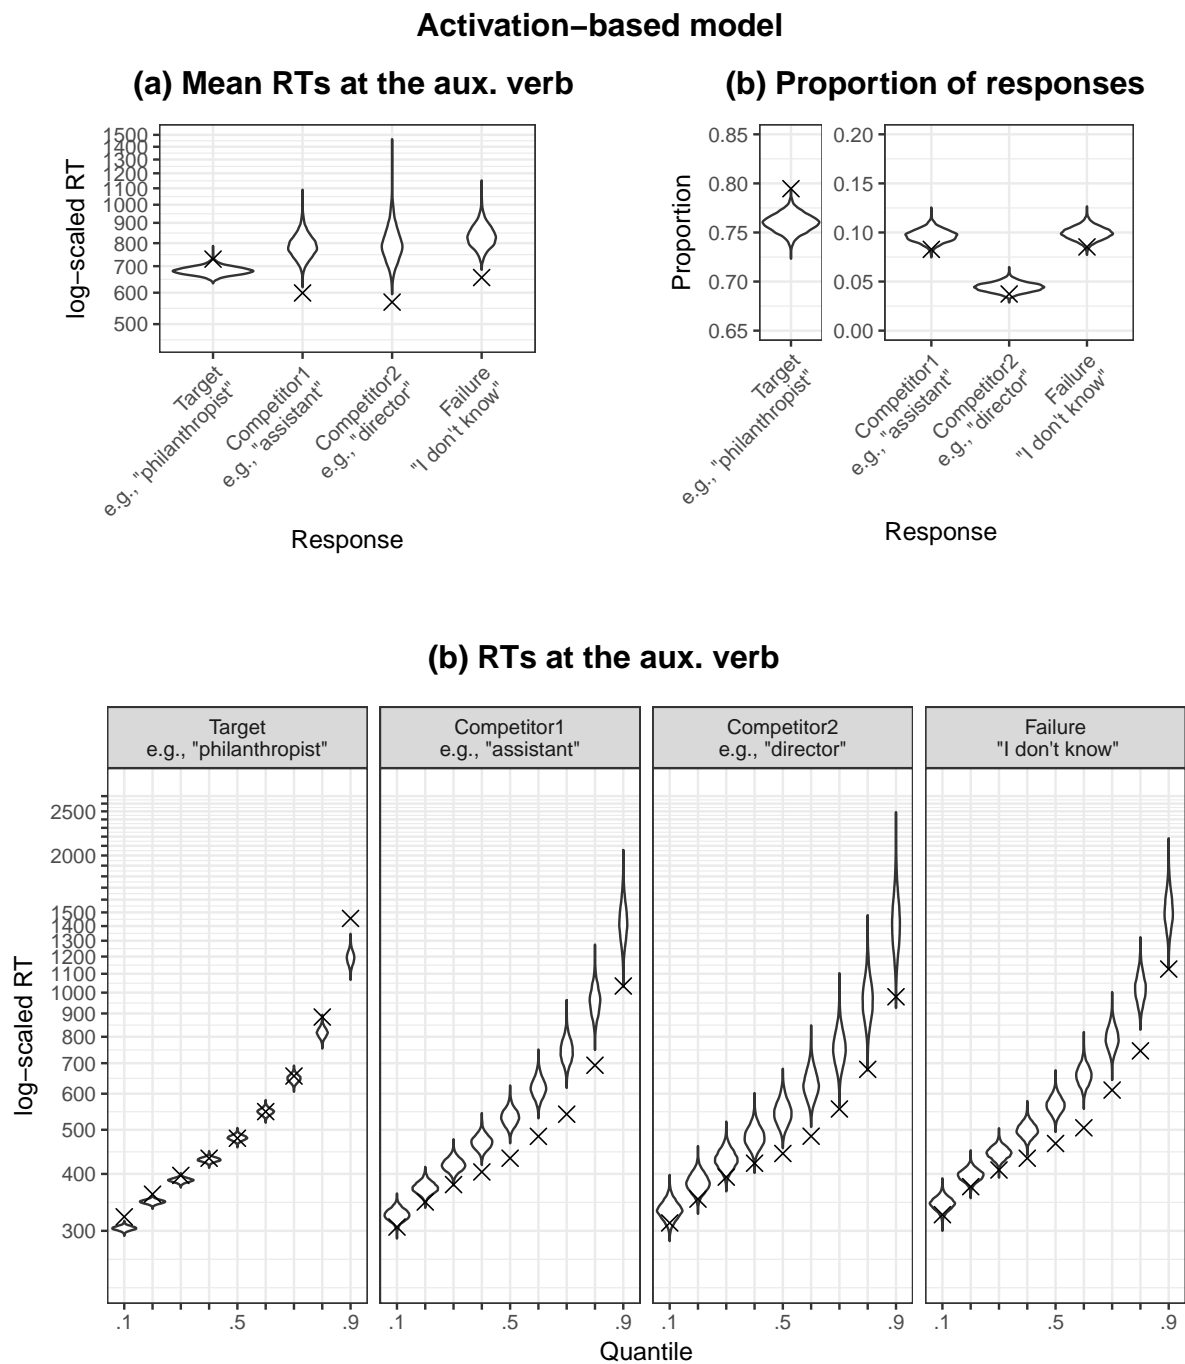


Figure 5. The top-most figure shows the fit of the mean reading times (RTs) for response (a) and proportion of responses (b) of the activation-based model. The width of the violin plot represents to the density of predicted mean RTs (a) and responses (b) generated by the model. The bottom figure (c) shows the fit of the .1–.9 quantiles of the reading times (RTs) for response of the activation-based model. The width of the violin plot represents to the density of predicted quantile generated by the model. The observed means and quantiles are represented with a cross.

accumulation of evidence for each choice assuming an arbitrary threshold of 10. In the activation-based model, these parameters represent the mean activation (together with a common additive constant) of the target, competitor NPs, and in the case of the failure option, the activation represents the speed of the timeout. As assumed by the activation-based model, the activation of the target is higher than the activation of the competitors and of the failure. The figure shows that the activations of the chunks fit the Equation (18), which indicates that the correct chunk should receive more activation on average than the competitor chunks.

Figure 6(b) shows the difference in activation for high and low interference for the target and the competitors. The figure shows that the posterior means and most of the probability mass of the effect of interference fit the predictions of Equations (17), i.e., interference lowers the activation of the target, and (19), i.e., interference raises the activation of the competitors. However, all the 95% CrIs include 0 as a plausible value; for the target: $\hat{\beta} = -0.03$, 95% CrI = $[-0.12, 0.06]$, $P(\beta > 0) = 0.25$, for the first competitor: $\hat{\beta} = 0.18$, 95% CrI = $[-0.06, 0.44]$, $P(\beta > 0) = 0.93$, and for the second competitor: $\hat{\beta} = 0.07$, 95% CrI = $[-0.27, 0.38]$, $P(\beta > 0) = 0.66$. There is clearly a lot of uncertainty about the parameters' plausible values. While the posteriors do not clearly contradict the predictions of the model regarding the effect of interference, they do not bring much support either.

Direct-access model.

Posterior predictive check. The posterior predictive check reveals that, in contrast to the activation-based model, the direct-access model is able to predict the main characteristics of the data fairly well. Figure 7(a) shows that the model is able to predict that reading times associated with correct responses are on average slower than the ones associated with incorrect ones, while Figure 7(b) shows that the model is able to predict fairly well the proportion of responses from the data. Furthermore, Figure 7(c) reveals that the fit is generally good for the entire distribution of reading times.




Figure 6. Mean activation of the different chunks assuming an arbitrary threshold of 10 (a), and mean difference between the activations of the chunks in high interference vs. low interference conditions (b). The outer error bars indicate 95% credible intervals while the inner error bars indicate 80% credible intervals.

Estimation of relevant parameters. The key parameters of the direct-access model are: (i) the probability that each of the candidate NPs would be retrieved (as shown in Figure 8), (ii) the probability of backtracking (reported below), and (iii) the time needed for backtracking (reported below). Figure 8(a) shows caterpillar plots of the posterior distributions for the parameters that represent probability of retrieving each chunk from memory in order to build a dependency at the auxiliary verb. Figure 8(a) shows that the retrieval of the target is more likely than the retrieval of the competitors or the retrieval failure; this is in agreement with Equation (21). Notice that since the model assumes that

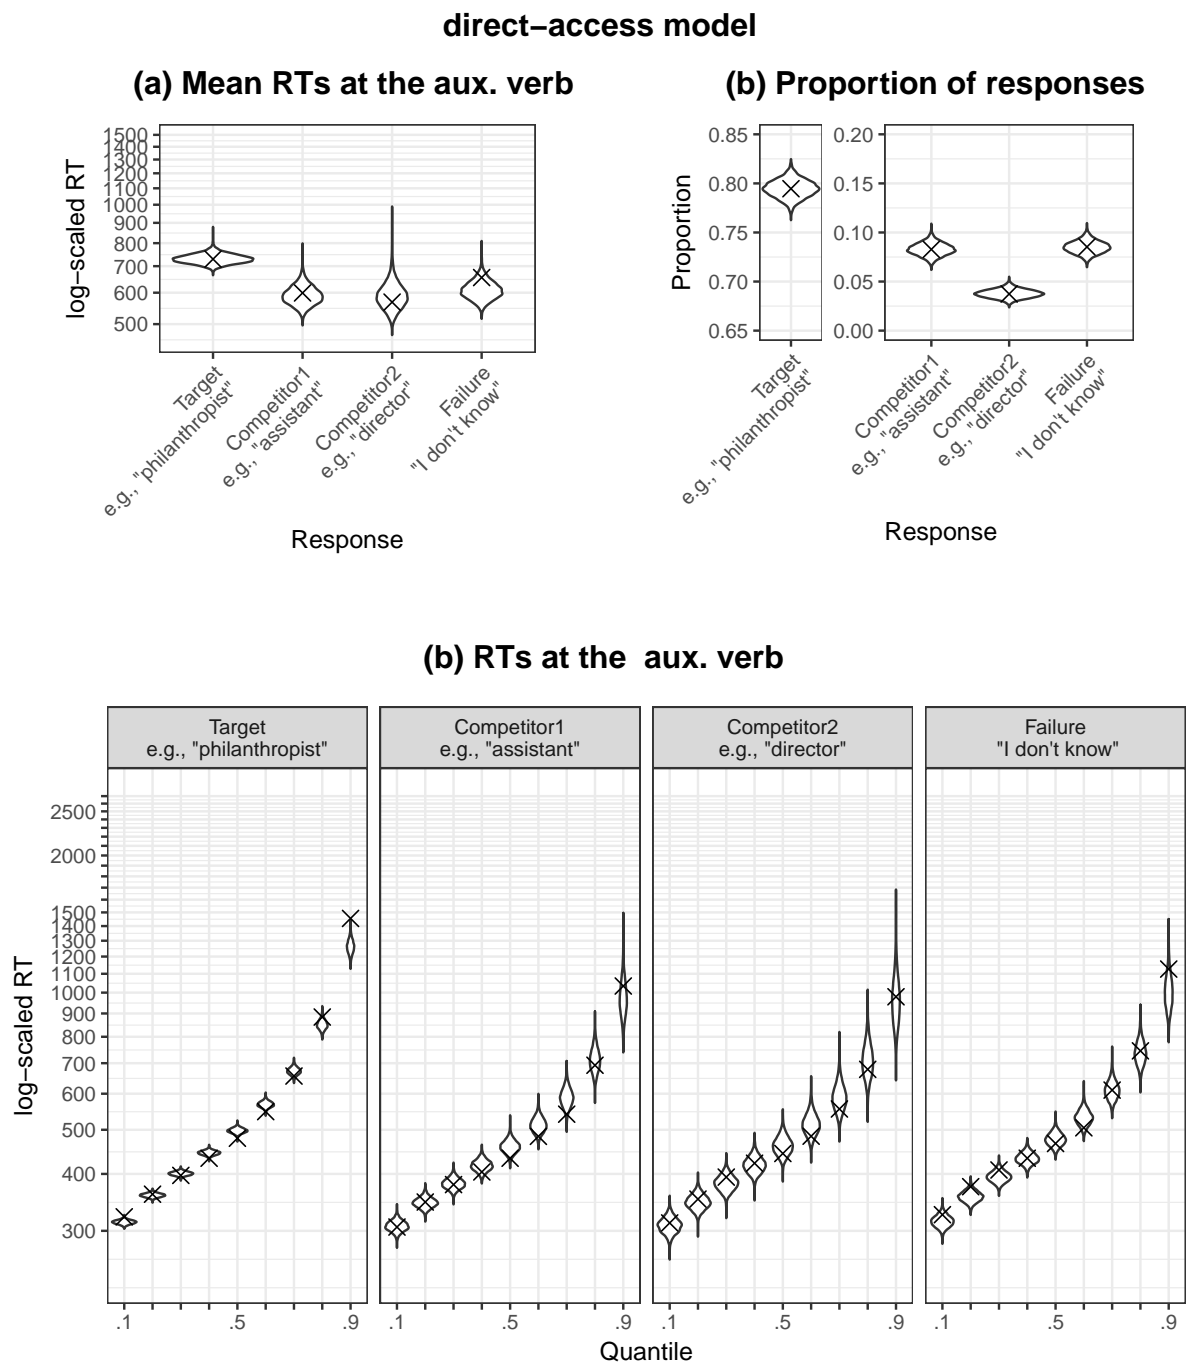


Figure 7. The top-most figure shows the fit of the mean reading times (RTs) for response (a) and proportion of responses (b) of the direct-access model. The width of the violin plot represents to the density of predicted mean RTs (a) and responses (b) generated by the model. The bottom figure (c) shows the fit of the .1–.9 quantiles of the reading times (RTs) for response of the direct-access model. The width of the violin plot represents to the density of predicted quantile generated by the model. The observed means and quantiles are represented with a cross.

backtracking is possible, after some trials the incorrect retrieval will be repaired. This means that the probability of retrieving a dependent is not the same as the proportion of times a response was given in the multiple choice task. In fact, the model estimates that around half of the time that there is a misretrieval, it will be corrected ($\hat{\theta}_b = 0.48$, 95% CrI = $[0.4, 0.55]$). In addition, the model estimates that backtracking takes 120 ms, 95% CrI = $[30, 231]$ ms (after transforming it from log-scale).

As shown in Figure 8(b) and similarly to the case of the activation-based model, the posterior distributions for the effect of interference do not clearly support or contradict the predictions of the model regarding the effect of interference. As before, the means and most of the probability mass of the posteriors for the effect of interference on the target and competitors fit the predictions of Equations (20), i.e., interference lowers the probability of retrieval of the target, and (22), i.e., interference raises the probability of retrieval of the competitors. However, all the 95% CrIs include 0; for the target: $\hat{\beta} = -0.04$, 95% CrI = $[-0.11, 0.02]$, $P(\beta > 0) = 0.08$, for the first competitor: $\hat{\beta} = 0.04$, 95% CrI = $[0, 0.08]$, $P(\beta > 0) = 0.96$, and for the second competitor: $\hat{\beta} = 0.02$, 95% CrI = $[-0.01, 0.04]$, $P(\beta > 0) = 0.88$.

Cross-validation: activation-based vs. direct-access models. In order to assess the compatibility of the models with the data, we compared how the models would generalize to an independent data set, that is, the pointwise out-of-sample prediction accuracy or \widehat{elpd} of the models. Comparing the models using k-fold cross validation reveals an estimated difference in \widehat{elpd} of -110 ($SE = 28$) in favor of the direct-access model ($\widehat{elpd} = -26747$, $SE = 100$) in comparison with the activation-based model ($\widehat{elpd} = -26858$, $SE = 100$).

Figure 9 shows for any given observation, whether one model has an advantage over the other in its predictive accuracy. Since higher (or less negative) values of \widehat{elpd} indicate a better fit for a model, observations that are further away from the dotted line correspond to data that are particularly better predicted by one model (and poorly by the other). This

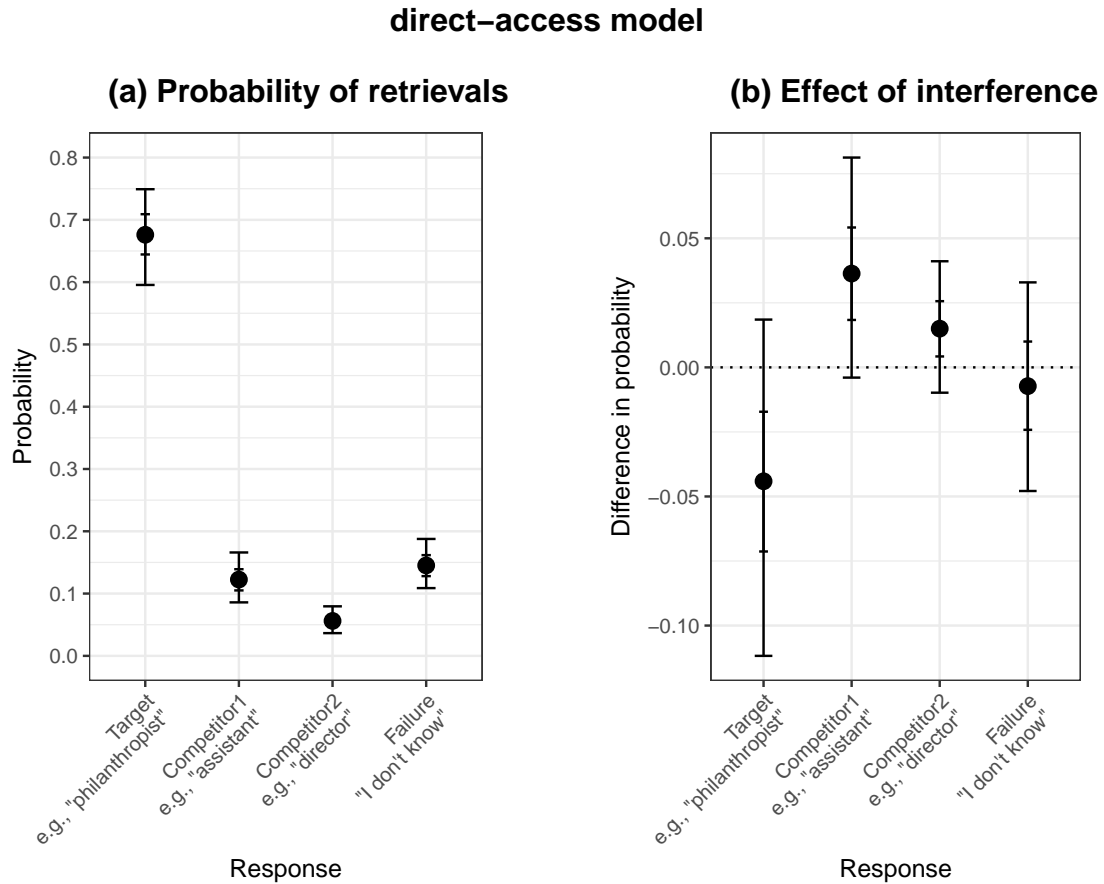


Figure 8. Mean probability of retrieval of the different chunks (a), and mean difference between the probabilities due to interference (b). The outer error bars indicate 95% credible intervals while the inner error bars indicate 80% credible intervals.

figure shows that the advantage of the direct-access model is not due to some outlier observations, but mostly due to a high number of observations that fit slightly better under this model than under the activation-based one (this is the darker patch on the top right corner). Figure 10 shows the difference between the \widehat{elpd} of the two models for every observation corresponding to either a correct or an incorrect response. The figure shows that most of the advantage of the direct-access model comes from reading times between 300 and 1000 ms (notice the darker patch above the zero dotted line). In addition, the direct-access model has a clear advantage in predicting long reading times associated with correct responses and short reading times associated with incorrect ones, while the

activation-based model has an advantage in predicting short reading times for correct responses and long reading times for incorrect ones.




Figure 9. Comparison of the activation-based and direct-access models in terms of their predictive accuracy for each observation. Each axis shows the expected pointwise contributions to 10-fold cross-validation for each model (\widehat{elpd} stands for the expected log pointwise predictive density of each observation). Higher (or less negative) values of \widehat{elpd} indicate a better fit. Darker cells represent a higher concentration of observations with a given fit.




Figure 10. Comparison of the activation-based and direct-access models in terms of their predictive accuracy for each observation depending on its log-transformed reading time (x-axis) and accuracy (left panel showing correct responses, and the right panel showing any of the possible incorrect responses). The y-axis shows the difference between the expected pointwise contributions to 10-fold cross-validation for each model (\widehat{elpd} stands for the expected log pointwise predictive density of each observation); that is, positive values represent an advantage for the direct-access model while negative values represent an advantage for the activation-based model. Darker cells represent a higher concentration of observations with a given fit.

Discussion

The evaluation of the activation-based and direct-access models reveals two sets of findings: one relates to the effect of interference on the key parameters of the models, and other the relates to their validity as models of retrieval in sentence comprehension.

Regarding the effect of interference on the parameters of the models, the results show that interference affects the parameters as expected, but some of the posteriors include a large degree of uncertainty. Given the relatively small effect of interference in reading times in the experimental study of Nicenboim et al. (Submitted), and given that we used a subset of the original data, this is not surprising. However, this serves as a sanity check that confirms that both models can work in principle and that experimental findings can produce the expected effects on the parameters of the models.

For the activation-based model, the underlying activation of the target NP was, as expected, clearly larger on average than the one of the competitors and the one associated with the timeout. The parameters that correspond to the effect of interference on activation, however, provided very weak evidence that interference decreases the activation of the target and increases the activation of the competitors.

Similarly for the direct-access model, the underlying probability of retrieving the target was clearly larger than the one of the competitors and the one associated with the failure of the retrieval process. The model estimated that approximately half of the time ($\hat{\theta}_b = 0.48$, 95% CrI = $[0.4, 0.55]$) that a retrieval was incorrect, it was repaired to the correct retrieval in 120 ms, 95% CrI = $[30, 231]$ ms. This finding shows, as McElree (2000) suggested, that it is possible to account for differences in reading times that arise only from differences in probabilities of retrieval, if there is a repair process that takes more than a negligible amount of time. However, as with the activation-based model, the posterior distributions present only weak evidence that interference decreases the probability of retrieving the target and increases the probability of retrieving one of the competitors.

In order to evaluate the validity of the models for retrieval in sentence comprehension, we examined whether the models were able to fit the patterns found in the data using posterior predictive checks, and we compared their predictive accuracy using cross-validation. The posterior predictive checks of the activation-based and direct-access models show clearly that some aspects of the data fit better under the direct-access model

than under the activation-based model. While the reading times at the auxiliary verb associated with correct responses in the multiple choice task were on average slower than the reading times associated with incorrect responses, this pattern could only be captured by the direct-access model. This is so because in the case of the direct-access model, reading times associated with correct responses are assumed to be a mixture of fast direct accessed retrievals and slower backtracked responses, while incorrect responses are assumed to be just direct accessed wrong or failed retrievals. By contrast, in the case of the activation-based model, reading times associated with correct responses are assumed to be faster on average than reading times associated with incorrect responses. The activation-based model assumes a race between the accumulation of evidence for the candidates to the retrieval, where the fastest item is the one retrieved. The particular characteristics of this race-between-accumulators model, which are motivated by ACT-R, include the assumption of a ballistic race (lack of fluctuations occurring during the accumulation process or within-choice noise; Brown & Heathcote, 2005), the same variance parameter for all the accumulators (i.e., a single between-choice noise), and no variation in the initial bias (Heathcote & Love, 2012). Under this type of race, the correct responses, which are answered more frequently than the incorrect ones, will also be the fastest on average.

Even though reading times for correct responses were on average slower than the ones for incorrect ones, this was not the case for every observation. Model comparison using cross-validation shows that the advantage of the direct-access model is based mainly on giving a better fit for reading times between 300 and 1000 ms, while the model is worst suited to predict fast reading times corresponding to correct responses and slow reading times corresponding to incorrect responses, which are better predicted by the activation-based model (see Figure 10).

Cross-validation supports the direct-access model, but does not rule out others models that assume a race between accumulators of evidence for each retrieval candidate:

the concept of activation determining retrieval latencies and accuracy may still be fruitful. It may be possible to explain the pattern in the data by including a mixture process in the race model, that is, if it is assumed, in a similar way as with the direct-access model, that the reading times associated with the correct responses are a mixture of fast retrievals due to high activation together with repaired wrong or failed retrievals. However, a model like this would be too flexible for the data at hand and may present problems of identifiability (since it would be hard to estimate the activation of the non-retrieved candidates).

A closely related model that has been proposed to account for fast errors by Nicenboim et al. (2016) assumes that failed retrievals may take less time than completed retrieval. This is achieved by assuming that, when the activation is too low, the retrieval is aborted instead of waiting until the timeout is reached. However, this would only explain the fast failures (“I don’t know answers” in the multiple choice task), but it would still leave fast retrievals of competitor NPs unexplained.

As we mentioned before, the activation-based model is based on a specific race model, namely the lognormal race model, which in turns is a very specific implementation of a model that assumes the sequential sampling of evidence for a decision (a class of models that includes the race of accumulators and random walk/diffusion models; for a review, see Ratcliff, Smith, Brown, & McKoon, 2016). There are other tasks that trigger error responses that are on average faster than correct responses and have been explained with sequential sampling models such as the drift diffusion model (Ratcliff & Rouder, 1998; Wagenmakers, Ratcliff, Gomez, & McKoon, 2008), or the linear ballistic accumulator (Brown & Heathcote, 2008). These models, in contrast with the lognormal race model, allow for variation in the initial bias due to noise. Fast errors can be captured by assuming a low threshold of evidence for all the decisions (in comparison with the variation in the initial bias). The low thresholds will produce faster responses in general, but the increase in speed will be larger for error responses. This is because incorrect responses will tend to occur when the process starts near the error boundary for the drift diffusion model or near

the thresholds of evidence of the incorrect responses for the linear ballistic accumulator (Wagenmakers et al., 2008; Heathcote & Love, 2012). This is in contrast with the situation where the thresholds of evidence are higher in comparison with the initial bias, in this case errors will tend to occur when the target candidate was slower than average. Since models such as the drift diffusion model or the linear ballistic accumulator allow fast or slow errors depending on the participants' decision thresholds, an interesting implication of the fast errors in the current experiment would be that participants privileged speed at the expense of accuracy without explicit instructions. Even though the aforementioned models could account for the faster average reading times associated with incorrect responses, they would lose the close connection with the ACT-R framework that motivated our use of the lognormal race model.

Regarding the lognormal race model, its limitation is that if equal variance is associated with each accumulator, fast errors on average cannot be predicted because bias (distance) and rate of accumulation cannot be disentangled (Heathcote & Love, 2012; Rouder et al., 2014). Fast errors on average can be predicted, if the variance of the accumulators of the incorrect responses is larger than the one of the correct response. Heathcote and Love (2012) propose that poorer matches may spread not only weaker activation on average but may also be noisier than stronger matches. This idea is also present in SAM framework of Gillund and Shiffrin (1984), which assumes slower reactions and more variability for poorer matches than for more precise matches. Figure 11 shows graphically how the distribution of correct and incorrect retrievals is generated from the activation-based model with different variances.

In the following section, we evaluate the activation-based model with different variances, one for the accumulator of target and one for the other accumulators, and we compare it with the direct-access model.

Evaluation of the activation-based model with different variances

We evaluated the activation-based model with different variances using the same data as with the previous models. As before, we (i) examined the descriptive adequacy of the model using posterior predictive checks, (ii) estimated its relevant parameters, and (iii) compared it with the direct-access model using cross-validation.

The assumptions of the activation-based model with different variances are identical to those of the default activation-based model, except that the noise in the rate of accumulation of evidence of each chunk can have different variances. This means that the lognormal distributions associated with each activation have different scale parameters (which corresponds to the standard deviation of the associated normal distribution). Since all the competitors were retrieved only 21% of the time, for simplicity (and for improving the convergence of the models) we assumed only two variances, one for the lognormal distribution associated with the target, and one for the competitors or the failure timeout.

Results

Posterior predictive check. The posterior predictive check reveals that the activation-based model with different variances can capture the main characteristics of the data. Figure 12(a) shows that the model predicts a wide range of reading times associated with the incorrect responses, and most of the predicted reading times associated with incorrect responses are only slightly faster than the correct ones. Figure 12(b) shows that the model is able to predict the proportion of responses from the data. Figure 12(c) reveals that the fit is better for the first half of the quantiles, while for the second half of the quantiles the predicted data contains the observed quantiles, mainly because of the wide distribution of predicted reading times.

Estimation of relevant parameters. The estimation of the key parameters and the relationship between parameters of the activation-based model with different variances shows similar results to the ones in the default activation-based model. As in the default

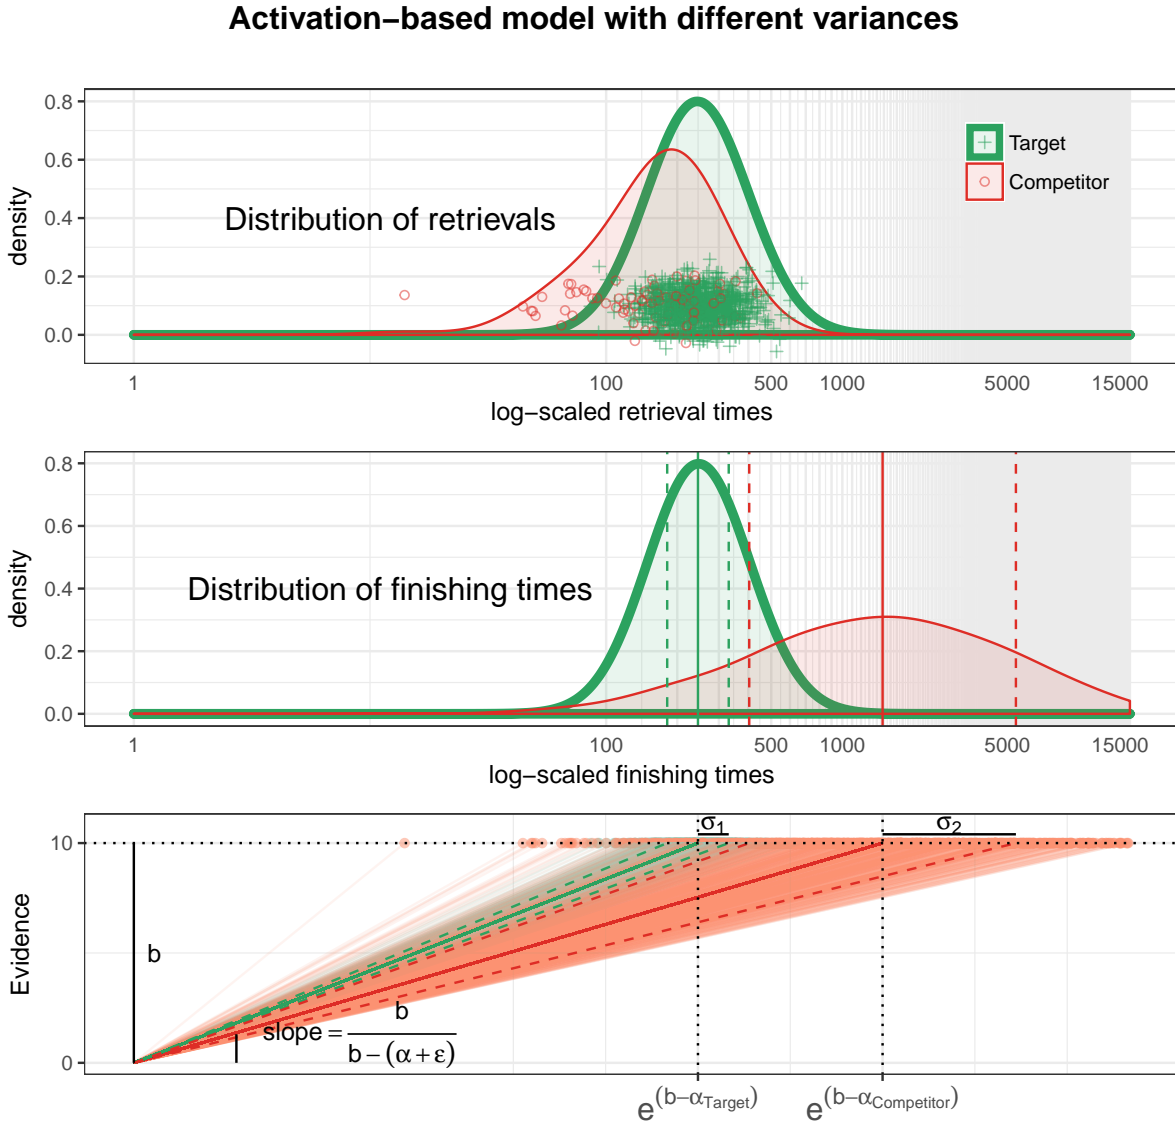


Figure 11. The figure depicts how the distribution of retrievals is generated from the activation-based model with different variances. The bottom figure depicts the parameters of the activation-based model, the full lines in green and red are the mean finishing times $t_{c=\text{Target}}$ and $t_{c=\text{Competitor}}$, and the broken lines are finishing times one standard deviation away from the mean. The middle figure shows the distributions of finishing times for target and candidate; since every chunk is associated a potential finishing time, t , both distributions have the same number of elements. The top-most figure shows the distribution of retrieval times (adding the shift parameter ψ would transform it to reading times); since only the winning chunks are retrieved, the distribution of retrieval times for targets has more elements than the distribution of the competitors. Notice that even though the finishing times for the competitors are slower on average than those of the targets (middle plot), the situation is reversed for the retrieval times (top-most plot).

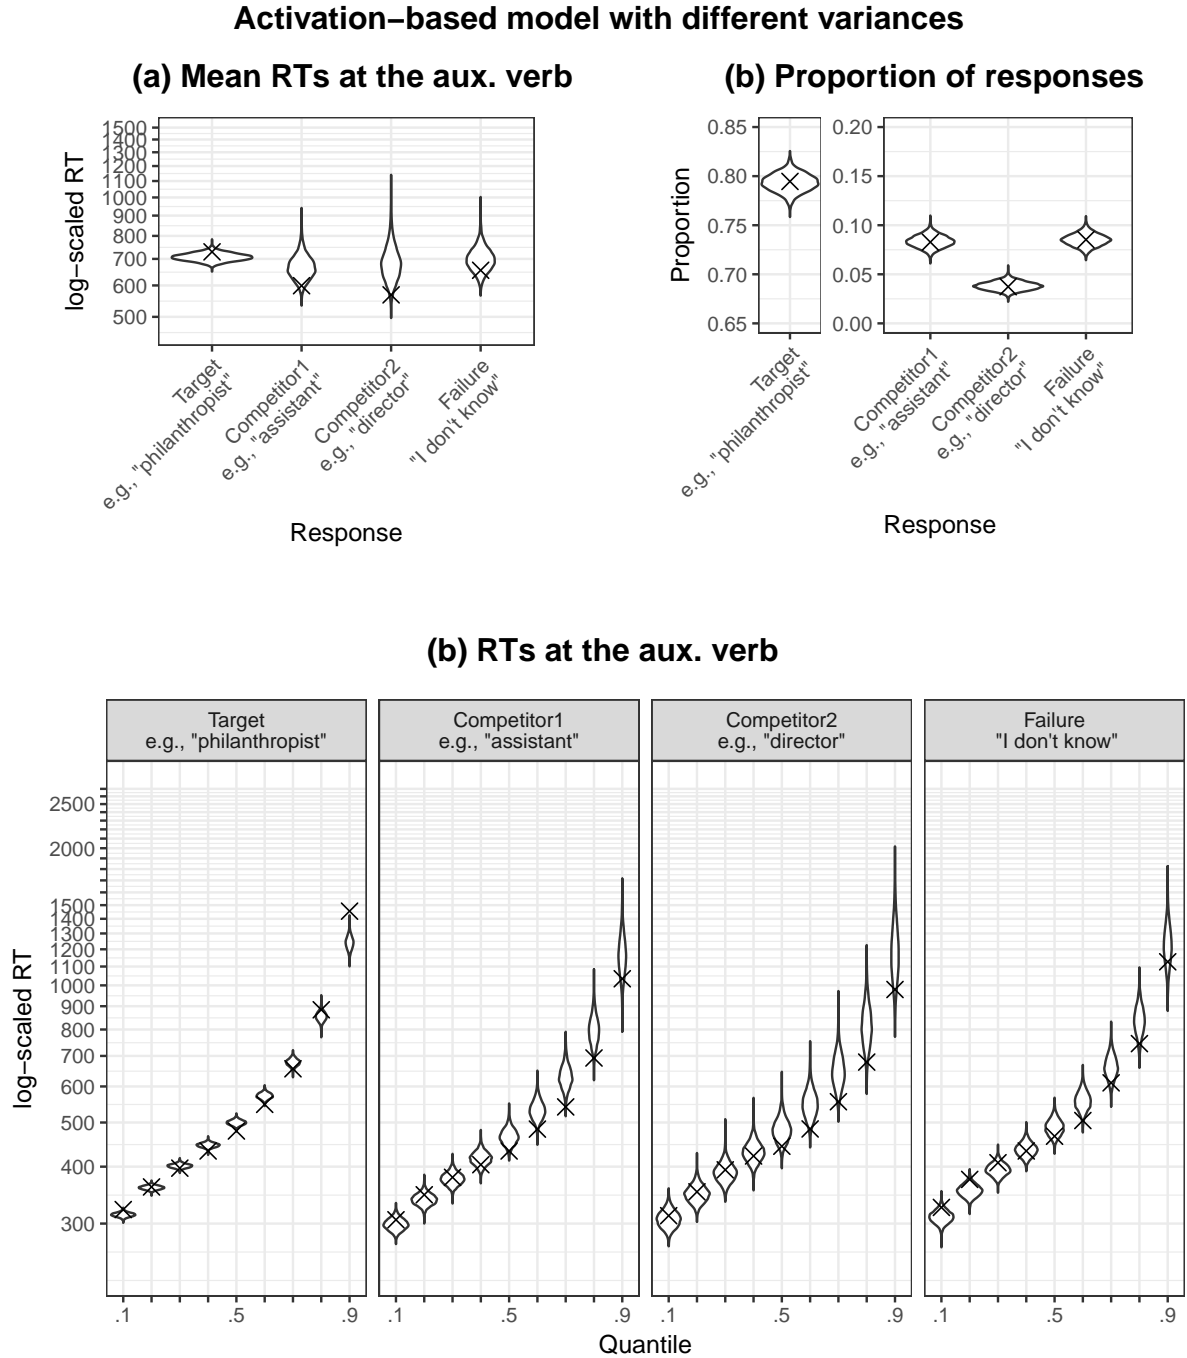


Figure 12. The top-most figure shows the fit of the mean reading times (RTs) for response (a) and proportion of responses (b) of the activation-based model with different variances. The width of the violin plot represents to the density of predicted mean RTs (a) and responses (b) generated by the model. The bottom figure (c) shows the fit of the .1-.9 quantiles of the reading times (RTs) for response of the activation-based model with different variances. The width of the violin plot represents to the density of predicted quantile generated by the model. The observed means and quantiles are represented with a cross.

model, Figure 13(a) shows that the activation of the target is higher than the activation of the competitors and of the failure. In addition and also similarly to the case of the default model, Figure 13(b) shows that while the posterior means and most of the probability mass of the effect of interference on the target and competitors fit the predictions of Equations (17) and (19), all the 95% CrIs include 0; for the target: $\hat{\beta} = -0.03$, 95% CrI = $[-0.11, 0.05]$, $P(\beta > 0) = 0.23$, for the first competitor: $\hat{\beta} = 0.17$, 95% CrI = $[-0.09, 0.46]$, $P(\beta > 0) = 0.89$, and for the second competitor: $\hat{\beta} = 0.13$, 95% CrI = $[-0.23, 0.49]$, $P(\beta > 0) = 0.77$. Recall that the variance was allowed to be different for the correct and incorrect retrievals; Figure 14 shows that, as hypothesized, this allows the scale associated with the distribution of activations of the incorrect retrievals to be larger than the one associated with correct retrievals.

Cross-validation: activation-based model with different variances vs. direct-access model. A comparison of the activation-based model with different variances and direct-access model using 10-fold cross-validation shows that the estimates of \widehat{elpd} are very similar, with a very small advantage for the activation-based model with different variances ($\widehat{elpd} = -26725$, $SE = 98$) in comparison with the direct-access model ($\widehat{elpd} = -26747$, $SE = 100$), namely an estimated difference in \widehat{elpd} of -22 ($SE = 18$); while the advantage of the direct-access model in comparison with the default activation-based model was of -110 with $SE = 28$. In addition, the activation-based model with different variances shows a clear advantage in comparison with the default activation model, namely an estimated difference in \widehat{elpd} of -133 ($SE = 18$)

Figure 15 shows that the predictive accuracy of the models is fairly similar with most of the observations being fit well by both of them. There are, however, some observations scattered at the bottom left corner of Figure 15, which favors the activation-based model with different variances. Figure 16 shows in blue cells the difference between the \widehat{elpd} of both models for every observation corresponding to either a correct or an incorrect response; and in white cells the previous comparison (from Figure 10) of the default

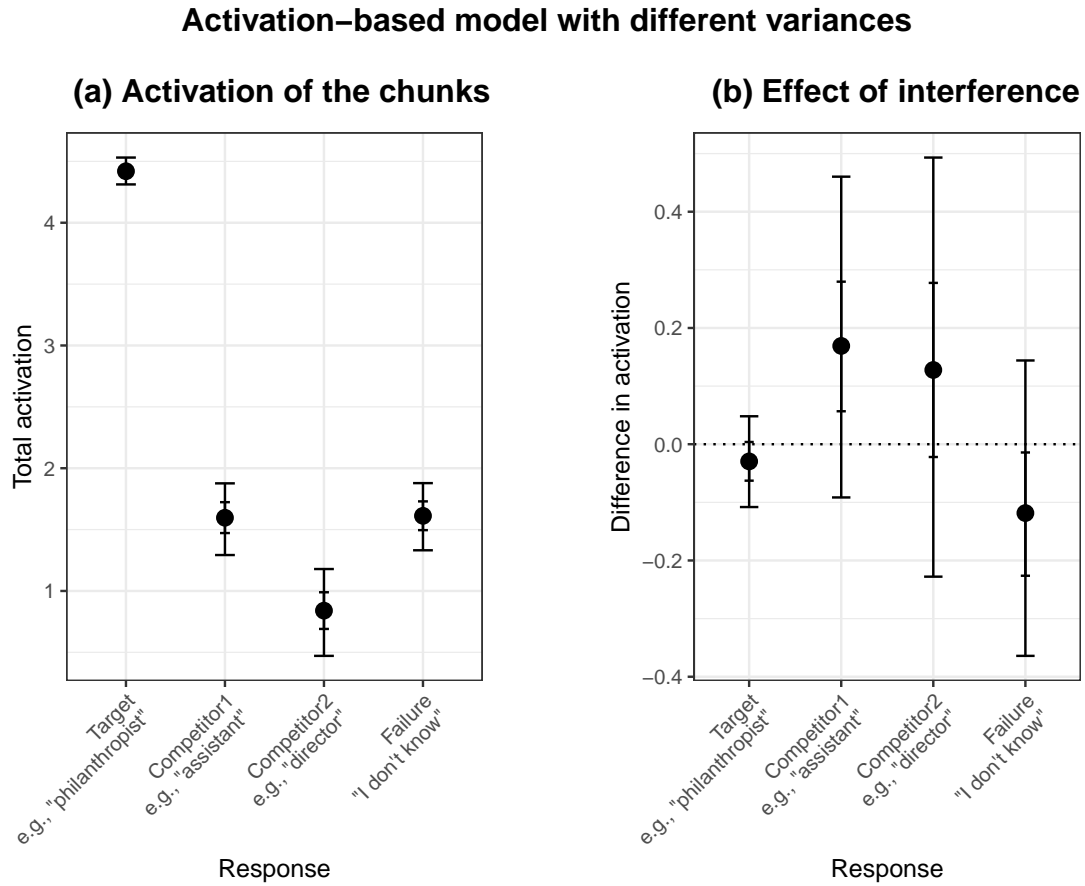


Figure 13. Mean activation of the different chunks assuming an arbitrary threshold of 10 (a), and mean difference between the activations due to interference (b). The outer error bars indicate 95% credible intervals while the inner error bars indicate 80% credible intervals.

activation-based model and direct-access model. Figure 16 shows that the difference between the activation-based model with different variances and direct-access model is smaller than the difference between the direct-access model and the (default) single-variance activation-based model. The main difference between the fits is that the activation-based model with different variances is able to account better for some fast and slow reading times associated with incorrect responses.




Figure 14. The figure depicts that the scale of the distributions of activations of the target chunk and of the competitors or timeout. The outer error bars indicate 95% credible intervals while the inner error bars indicate 80% credible intervals.

Discussion

The estimation of the relevant parameters of the activation-based model with different variances shows that the scale parameter associated with the distribution of activations for incorrect retrievals is larger than the one associated with correct ones, as it is necessary to account for fast errors. However, this did not change the predicted interference effect compared to the default activation-based model. Similarly, as with the default model, the parameters that correspond to the effect of interference on activation do not clearly support or contradict the predictions of the model regarding the effect of interference, that is that interference decreases the activation of the target and increases the activation of the competitors.

Regarding the descriptive adequacy of the model, even though the inclusion of different variances improves the fit, the posterior predictive checks show more variation on the predicted reading times associated with incorrect responses for this model than for the




Figure 15. Comparison of the activation-based model with different variances and the direct-access model in terms of their predictive accuracy for each observation. Each axis shows the expected pointwise contributions to 10-fold cross validation for each model (\widehat{elpd} stands for the expected log pointwise predictive density of each observation). Higher (or less negative) values of \widehat{elpd} indicate a better fit. Darker cells represent a higher concentration of observations with a given fit.

direct-access model. This is not necessarily a disadvantage, and it may indicate that the direct-access model is more flexible and may be slightly overfitting the data, since these predictions are generated with the best estimates (and posterior distributions) to account for the data. In fact, despite an apparent better fit for the direct-access model, the estimates of predictive accuracy (\widehat{elpd}) are very similar with a very slight advantage for the activation-based model with different variances. In contrast with the difference between the fit of the original models (i.e., the default activation-based vs. the direct-access model shown in Figure 10), the difference between the fit of the activation-based model with different variances and the direct-access model is smaller (see Figure 16), with the new version of the activation-based model giving a better fit to some of the fast and slow reading times associated with incorrect responses.

This comparison shows that even though the inclusion of different variances for the




Figure 16. Comparison of the activation-based model with different variances and direct-access model in terms of their predictive accuracy for each observation depending on its log-transformed reading time (x-axis) or accuracy (left panel showing correct responses, and the second panel showing any of the possible incorrect responses). The y-axis shows the difference between the expected pointwise contributions to 10-fold cross-validation for each model (\widehat{elpd} stands for the expected log pointwise predictive density of each observation); that is, positive values represent an advantage for the direct-access model while negative values represent an advantage for the activation model with different variances. Darker cells represent a higher concentration of observations with a given fit. The white cells show the comparison (shown earlier in Figure 10) of the default activation-based model with the direct-access model.

accumulators does not imply a clear superiority over the direct-access model, it is possible to account for the data with a model which is based on a race of accumulation of evidence.

General discussion

We evaluated two models that have been successful in explaining similarity-based interference in sentence comprehension: Lewis and Vasishth’s (2005) activation-based model following ACT-R assumptions (Anderson et al., 2004) and McElree’s (2000) direct-access model. We also evaluated a third model, a variation of the activation-based model.

The models were implemented in a Bayesian hierarchical framework and fit to the data of Nicenboim et al. (Submitted). Even though the activation-based model was already implemented computationally (Lewis & Vasishth, 2005), our implementation enabled us to go beyond simulations as they are usually done for this model (e.g., Vasishth & Lewis, 2006; Nicenboim et al., 2016), and to fit the observations of an experiment by accounting for variation coming from participants and experimental items. For the direct-access model, we provided a first computational implementation which allowed us to derive precise and unambiguous predictions, which are fully transparent in our instantiation of the model. We first summarize our findings, and we then discuss the motivation of this work, the implications of the findings, and future work.

Our evaluation can be summarized in three main results. First, the underlying parameters of both models behave as expected under interference effects. However, the parameters showed a large degree of uncertainty in their posterior distributions. While this may be due to the relatively small magnitude of the interference effect in the original experiment (Nicenboim et al., Submitted), the findings confirm that, as expected, both models can in principle explain interference effects.

Second, we evaluated the validity of both models in predicting the reading times and accuracy patterns during retrieval. The posterior predictive checks and the comparison using cross-validation show that some aspects of the data fit better under the direct-access model than under the default activation-based model. The data showed on average slower reading times associated with correct responses than with incorrect ones, and this pattern

could be explained only by the direct-access model. This suggests that the default activation-based model may not be flexible enough to accommodate patterns in the data that go beyond means between conditions.

Third, we show that by introducing a modification to the default activation-based model, namely, by assuming that the accumulation of evidence for the retrieval of incorrect items is not only slower but noisier, the new model can provide a fit as good as that of the direct-access model.

The importance of a formal comparison of Lewis and Vasishth's (2005) activation-based model and McElree's (2000) direct-access model lies in disentangling their predictions. Since both models assume that dependencies of non-adjacent elements are created via a content-addressable cue-based retrieval mechanism, they have been used almost interchangeably to explain interference effects (e.g. Van Dyke & McElree, 2006). For experiments that draw inferences from differences in means, these two models yield identical predictions for the inhibitory effect of similarity-based interference: namely, longer reading times at the retrieval of a dependent and/or a reduction of comprehension accuracy when several items share a feature associated with a retrieval cue. However, these models are based on different underlying assumptions. The activation-based model follows ACT-R assumptions (Anderson et al., 2004); in this framework, the activation of the items in memory determines the retrieval accuracy and latency, and the activation of the target of retrieval is, in turn, adversely affected by interference. Crucially, latency and accuracy are not deterministic because activation fluctuates due to noise in the system. We show that this process can be seen as a lognormal race between accumulators of evidence with a single variance for all the accumulators, where activation represents the rate of accumulation of evidence. In contrast, the direct-access model assumes a model of memory where only the probability of retrieval can be affected by interference, while items take the same time to be retrieved (if they are not in the focus of attention as it is the case for non-local dependencies). In this model, differences in latencies are a by-product of the possibility of

backtracking and repairing incorrect retrievals. These different assumptions lead to a different behavior in the relationship between reading times and response accuracy on a trial-level basis which cannot be examined by only comparing mean reading times or accuracy between conditions. While acceptability judgment tasks with speed-accuracy trade-off (SAT) allow a finer grain look at the reaction times and have been used to argue in favor of the direct-access model (see, for example, Van Dyke & McElree, 2011), until now there has been no computational evaluation of the model in reading for comprehension.

While the activation-based model uses the declarative retrieval module of ACT-R, which has been shown to be an empirically successful model (e.g., Anderson et al., 1998; Anderson & Reder, 1999; Van Rijn & Anderson, 2003), our findings show that its default implementation cannot account for incorrect retrievals that were generally faster than the correct ones in our data. The model cannot account for this pattern because items in memory that match the retrieval cues will have higher activation on average than competitors that match the retrieval cues only partially. The higher activation on average leads in turn to overall faster retrievals. In contrast, the direct-access model can successfully accommodate faster incorrect retrievals. This is done by assuming that reading times associated with correct responses are generated from a mixture distribution of fast directly accessed correct retrievals at the first attempt together with slower backtracked and repaired retrievals. Reading times associated with incorrect responses, in contrast, belong to a faster distribution of retrieval latencies of items that are directly accessed. It should be noted that this repair mechanism that explains slow correct retrievals could in principle be added to the activation-based model, but it would lead to an unidentifiable model. The direct-access model, however, is able to account for the data with a very simple architecture that can integrate this repair mechanism.

While the simple architecture of the direct-access model may be preferred on grounds of parsimony, the activation-based model has some arguably desired characteristics: it is compatible with a sequential sampling framework (such as the drift diffusion model:

Ratcliff, 1978; the leaky competitive accumulator: Usher & McClelland, 2001; linear deterministic models: Heathcote & Love, 2012, among others). In the sequential sampling framework, decisions (such as which is the right dependent that needs to be retrieved) are considered a process of noisy accumulation of evidence, which has been shown to be compatible with the behavior of populations of neurons (e.g., Zandbelt, Purcell, Palmeri, Logan, & Schall, 2014). In addition, sequential sampling has been also linked to theories of optimality (Ratcliff et al., 2016; Summerfield & Tsetsos, 2015), which compare how an ideal agent would perform (given the levels of uncertainty in the stimuli) with the actual behavior of participants.

The sequential sampling framework could still be useful to explain retrieval, if we assume that the retrieval process behaves similarly to other more complex accumulator models such as the linear ballistic accumulator (Brown & Heathcote, 2008), since, as explained earlier, they can account for fast and slow errors (Wagenmakers et al., 2008). However, these models lose the close connection with the ACT-R framework that motivated the lognormal race, which underlies the activation-based model. In addition, given that these models are more complex than the lognormal race model, it may be the case that they overfit the data and it is not clear whether their fit would be comparable to the fit of the direct-access model. A potential future direction of this work would be to evaluate different plausible accumulator models as models of retrieval.

We also relaxed one of the assumptions of ACT-R to capture the patterns of the data: Here we assumed that the activation of chunks that match the retrieval cues only partially is not only lower but also noisier. This amounts into assuming different variances for the different accumulators. Heathcote and Love (2012) show that when the accumulators associated with incorrect responses have a larger variance than the accumulator of correct responses, the model can account for fast errors on average. For simplicity, we assumed one variance for the accumulator of the target, and one for the competitors and failure accumulators. While our study shows that this is enough to account for the pattern in the

data, nothing would prevent all accumulators from having different variances.

Both the activation-based model with different variances and the direct-access model showed an equally good fit to the data. In order to investigate these models' relative fit, future work should replicate the classical interference results (e.g., Van Dyke & McElree, 2006; Van Dyke, 2007; Van Dyke & McElree, 2011), while including reading times and questions probing the comprehension of the relevant dependencies. Given that the posteriors for the interference effect in the relevant parameters had a broad distribution and were not too informative, a large-sample replication of the classical interference experiments would allow us to verify the main predictions of the cue-based retrieval models.

There are other phenomena that the models could explain. These are: (i) the facilitatory interference effects found in ungrammatical sentences (Wagers, Lau, & Phillips, 2009), (ii) the ambiguity advantage in relative clauses (Traxler, Pickering, & Clifton, 1998) and the effect of task demands (Swets, Desmet, Clifton, & Ferreira, 2008), and (iii) good-enough processing (Ferreira et al., 2002). Since some of the predictions of the activation-based model with different variances are not very intuitive, we provide an *R* script called *race-plot* using the *Shiny* package (Chang, Cheng, Allaire, Xie, & McPherson, 2016) that can help visualizing the predictions.⁴

Facilitatory interference. Wagers et al. (2009) noticed that the so-called number attraction effect in ungrammatical sentences such as (4), that is, the speedup in *are* in (4b) vs. (4a), could be accounted by Lewis and Vasishth's (2005) activation-based model.

- (4) a. *The key_{+sing} to the cabinet_{+sing} are in the box.
 b. *The key_{+sing} to the cabinets_{+plur} are in the box.

In sentences like (4), a cue-based retrieval mechanism would assume that a retrieval is initiated at the verb (*are*) with at least two retrieval cues: grammatical subject and plural.

⁴The application can be accessed in the browser with the following commands in R:
`install.packages(c("dplyr", "tidyr", "ggplot2", "cowplot", "shiny")) #if needed`
`library(shiny) #load shiny`
`runUrl("http://www.ling.uni-potsdam.de/~nicenboim/code/race-plot.zip")`

In sentence (4a), *the key* matches one of the retrieval cues, because it is the grammatical subject, but mismatches the plural cue. In sentence (4b), both nouns partially match the retrieval cues: *the key* matches the grammatical subject cue, while *the cabinets* matches the plural cue. An interesting prediction of the activation-based model (confirmed by experimental findings; see Jäger et al., 2017, for a meta-analysis) is that reading times are *faster* at the verb in (4b) than in (4a). This is so because a situation with no unambiguous match (both nouns are partial matches) leads to statistical facilitation (Raab, 1962), that is, an overall speedup when we examine mean reading times (facilitatory interference). For facilitatory interference in ungrammatical sentences, the predictions of the default activation-based model and the activation-based model with different variances are the same. This situation can be simulated using *race-plot* script mentioned before, by assigning arbitrary (but plausible) activations to the candidates to retrieval in (4a) and (4b): In (4a), *the key* (partial match) can be assigned an activation of 4 and *the cabinet* (no match) an activation of 2.5 (and $\sigma = 1.5$); this would result in a mean reading time of approximately 832 ms. Notice that since the process is not deterministic, different simulations will show different retrieval times; the relationship between the conditions, however, should hold on average. In (4b), *the key* (partial match) can be assigned an activation of 4 and *the cabinets* (partial match) an activation of 3.5 (since they will not necessarily reach exactly the same activation); this would result in a faster reading time on average, approximately 692 ms.

In contrast to the activation-based model, the direct-access model would not predict a difference in reading times at the verb between (4a) and (4b). This is the case since increased reading times depend only on backtracking, which would only occur to repair an initially incorrect retrieval. In ungrammatical sentences with partial match, it is unclear how the repair would work, and why there would be more backtracking in (4a) than in (4b).

The predictions of the activation-based model, however, have not been investigated taking into account both reading times and comprehension. Even if a speedup compatible with facilitatory interference has been reported in the literature, the activation-based

model would be accounting for facilitatory interference only if participants reach a different interpretation of the sentence in (4b) more often than in (4a).⁵

The ambiguity advantage in relative clauses and task-demands effects.

The so-called ambiguity advantage is based on the observation of Traxler et al. (1998), who found a speedup at *mustache* in ambiguous conditions such as (5c) in comparison with unambiguous conditions such as (5a) and (5b), where *mustache* is the disambiguating word.

- (5) a. The driver of the car that had the mustache was pretty cool. (high attachment)
- b. The car of the driver that had the mustache was pretty cool. (low attachment)
- c. The son of the driver that had the mustache was pretty cool. (globally ambiguous)

The account of the activation-based model with different variances is very similar to the unrestricted race model proposed by van Gompel, Pickering, and Traxler (2000), which predicts statistical facilitation in the case of ambiguity. According to the unrestricted race model, the parser starts building all possible structures simultaneously. While the time taken depends on plausibility, it is also affected by noise. This means that the adopted structure in each trial is the one that takes the least time, leading to shorter time on average when there are more candidates.

The activation-based model with different variances would yield similar predictions to the unrestricted race model if the inhibitory effect of interference in (5c) is sufficiently small. Given the relatively small magnitude of the interference effect in the literature (Jäger et al., 2017; Nicenboim et al., Submitted), this is likely to be the case. In unambiguous cases such as (5a) and (5b), there is only one NP that matches the retrieval cue: “being capable of having a mustache” (i.e., *the driver*). In ambiguous cases such as (5c), there are two NPs that match the retrieval cue (i.e., *The son* and *the driver*).

Therefore, we would expect statistical facilitation (similarly to the case of facilitatory

⁵Notice that even though it is unlikely that readers would understand that *the cabinets are in the box*, it may be that the sentence is reanalyzed when the parser reaches *box*.

interference) leading to faster reading times on average. This situation can be simulated using the *race-plot* script similarly as before: In (5a) or (5b), *the driver* (full match) can be assigned an activation of 5 (and $\sigma = 1$), and *the car* (partial match) can be assigned an activation of 2.5 (and $\sigma = 2$); this would result in a mean reading time of approximately 416 ms. In (5c), both *The son* and *the driver* should have similar activation since there is no penalty component involved, both are a full match. However, the cue “can have a mustache” does not uniquely identify any candidate. Given the small magnitude of inhibitory interference effects, we could assume an activation of 4.8 (instead of 5) and the same variance since there is no mismatch (i.e. $\sigma = 1$) for both NPs. This would result in a faster reading time on average, approximately 349 ms.

Furthermore, the activation-based model may be able to account for Logačev and Vasishth’s (2016) observation that the parser seems to behave in a way that resembles a race between interpretations (low attachment vs. high attachment) but it is also task-dependent (as assumed by Swets et al., 2008). This could be achieved by setting the timeout (the parameters of the accumulator associated with the retrieval failure) to be task-dependent: longer timeouts when instructions or context encourage attentive reading and shorter timeouts when a full interpretation is not needed for successfully completing the experimental task.

In this case, the direct-access model could also predict the ambiguity advantage in a very simple way: While in (5a) or (5b) it is possible to retrieve the incorrect NP (i.e. “the car”) leading to a certain proportion of slower backtracked retrievals, in (5c) there should only be fast directly accessed retrievals, since both NPs (i.e. “The son” and “the driver”) are correct targets. In addition, for the direct-access model, the proportion of incorrect retrievals that are backtracked could be task dependent, with a larger proportion of backtracking associated with deeper processing. However, the predictions of both models would not be identical. The direct-access model predicts that the reading times at the disambiguating region when the incorrect interpretation (or no interpretation) is held in

(5a) or (5b) would be identical to the reading times in (5c). In contrast, for the activation based-model with different variances, the relationship between reading times at the different conditions would depend on comprehension accuracy. Future work that includes measures of reading times and queries for the comprehension of the relative clause, as well as manipulates task demands could compare the activation-based model with different variances, the direct-access model, and the model presented in Logačev and Vasishth (2016), which subsumes the unrestricted race model and allows it to be task-dependent.

Good-enough processing. While a comprehensive alternative to good-enough processing is out of the scope of this section (see Christianson, 2016, for a complete overview), it should be noticed that without further assumptions the activation-based model with different variances and the direct-access model can account for manipulations that show (sometimes unexpectedly) fast reading times which have been attributed to good-enough processing. For the activation-based model with different variances, this can be achieved by associating the timeout with either task demands as suggested previously or also with individual differences. An increase of either timeout speed (i.e., the rate of accumulation of the retrieval failure) or an increase of its noise (i.e., the variance of the accumulator associated with retrieval failure) would lead to more frequent shallow parses with incomplete dependencies which are read faster. Thus, experiments that probe the comprehension of certain dependencies less often may lead to faster (and maybe noisier) timeouts, which would in turn lead to shorter mean reading times. Individual differences in participants such as working memory capacity may have a similar effect, with lower capacity leading to faster and noisier timeouts in the retrieval process.

Similarly for the direct-access model, the probability of backtracking could be affected by task demands or by individual differences: A less demanding task would reduce reading times and comprehension accuracy on average by discouraging backtracking. Individual differences may have a similar effect, participants with lower working memory capacity may be less prone to backtracking. As suggested before, this could be assessed in

future work by including measures of reading times and comprehension accuracy of the relevant dependencies.

Conclusion

We have provided an evaluation of two theoretically grounded and empirically successful models in explaining similarity-based interference in sentence comprehension: Lewis and Vasishth's (2005) activation-based model built under the assumptions of ACT-R (Anderson et al., 2004) and McElree's (2000) direct-access model. We also evaluated a third model, a variation of the activation-based model.

Our evaluation, which consisted in implementing these models in a Bayesian hierarchical framework, confirms that, as expected, both the activation-based and direct-access models can in principle explain interference effects. However, posterior predictive checks and model comparison using cross-validation show that some aspects of the data fit better under the direct-access model, in particular, the default activation-based cannot predict that, on average, incorrect retrievals would be faster than correct ones.

Finally, we show that by introducing a modification of the activation model, namely, by assuming that the accumulation of evidence for the retrieval of incorrect items is not only slower but noisier (i.e., different variances for the correct and incorrect items), the new model can provide a fit as good as that of the direct-access model.

References

- Akaike, H. (1974). A new look at the statistical model identification. *IEEE Transactions on Automatic Control*, 19(6), 716–723.
- Anderson, J. R., Bothell, D., Lebiere, C., & Matessa, M. (1998). An integrated theory of list memory. *Journal of Memory and Language*, 38(4), 341–380.
doi:10.1006/jmla.1997.2553
- Anderson, J. R., Bothell, D., Byrne, M. D., Douglass, S., Lebiere, C., & Qin, Y. (2004). An integrated theory of the mind. *Psychological Review*, 111(4), 1036–1060.
doi:10.1037/0033-295x.111.4.1036
- Anderson, J. R. & Reder, L. M. (1999). The fan effect: New results and new theories. *Journal of Experimental Psychology*, 128(2), 186. doi:10.1037/0096-3445.128.2.186
- Audley, R. J. & Pike, A. R. (1965). Some alternative stochastic models of choice. *British Journal of Mathematical and Statistical Psychology*, 18(2), 207–225.
doi:10.1111/j.2044-8317.1965.tb00342.x
- Brown, S. D. & Heathcote, A. (2005). A ballistic model of choice response time. *Psychological Review*, 112(1), 117–128. doi:10.1037/0033-295X.112.1.117
- Brown, S. D. & Heathcote, A. (2008). The simplest complete model of choice response time: linear ballistic accumulation. *Cognitive psychology*, 57(3), 153–178.
doi:10.1016/j.cogpsych.2007.12.002
- Chang, W., Cheng, J., Allaire, J., Xie, Y., & McPherson, J. (2016). *Shiny: Web application framework for R*. R package version 0.14.
- Christianson, K. (2016). When language comprehension goes wrong for the right reasons: Good-enough, underspecified, or shallow language processing. *The Quarterly Journal of Experimental Psychology*, 69(5), 817–828. doi:10.1080/17470218.2015.1134603
- Cowan, N. (1995). *Attention and memory: An integrated framework*. Oxford University Press.

- Dillon, B., Chow, W. y., Wagers, M., Guo, T., Liu, F., & Phillips, C. (2014). The structure-sensitivity of memory access: Evidence from Mandarin Chinese. *Frontiers in Psychology*, 5(1025). doi:10.3389/fpsyg.2014.01025
- Dillon, B., Mishler, A., Sloggett, S., & Phillips, C. (2013). Contrasting intrusion profiles for agreement and anaphora: Experimental and modeling evidence. *Journal of Memory and Language*, 69(2), 85–103. doi:10.1016/j.jml.2013.04.003
- Engelmann, F. (2015). *Toward an integrated model of sentence processing in reading* (Doctoral dissertation, University of Potsdam).
- Engelmann, F., Jäger, L. A., & Vasishth, S. (Submitted). *The effect of prominence and cue association in retrieval processes: A computational account*. doi:10.17605/OSF.IO/B56QV
- Feldman, J. (2016). What are the “true” statistics of the environment? *Cognitive Science*, 1–33. doi:10.1111/cogs.12444
- Ferreira, F., Bailey, K. G. D., & Ferraro, V. (2002). Good-enough representations in language comprehension. *Current Directions in Psychological Science*, 11(1), 11–15. doi:10.1111/1467-8721.00158
- Foraker, S. & McElree, B. (2011). Comprehension of linguistic dependencies: Speed–accuracy tradeoff evidence for direct–access retrieval from memory. *Language and Linguistics Compass*, 5(11), 764–783. doi:10.1111/j.1749-818X.2011.00313.x.Comprehension
- Gelman, A., Carlin, J., Stern, H., Dunson, D., Vehtari, A., & Rubin, D. (2014). Bayesian data analysis. (Third). Taylor & Francis.
- Gelman, A. & Hill, J. (2007). *Data analysis using regression and multilevel/hierarchical models*. Cambridge, UK: Cambridge University Press.
- Gelman, A., Hill, J., & Yajima, M. (2012). Why we (usually) don’t have to worry about multiple comparisons. *Journal of Research on Educational Effectiveness*, 5(2), 189–211. doi:10.1080/19345747.2011.618213

- Gelman, A., Hwang, J., & Vehtari, A. (2014). Understanding predictive information criteria for Bayesian models. *Statistics and Computing*, 24(6), 997–1016.
doi:10.1007/s11222-013-9416-2
- Gibson, E. (2000). The dependency locality theory: A distance-based theory of linguistic complexity. In *Image, language, brain: Papers from the first mind articulation project symposium* (pp. 95–126).
- Gillund, G. & Shiffrin, R. M. (1984). A retrieval model for both recognition and recall. *Psychological Review*, 91(1), 1. doi:10.1037/0033-295X.91.1.1
- Gordon, P. C., Hendrick, R., & Levine, W. H. (2002). Memory-load interference in syntactic processing. *Psychological Science*, 13(5), 425–430. doi:10.1111/1467-9280.00475
- Grodner, D. & Gibson, E. (2005). Consequences of the serial nature of linguistic input for sentential complexity. *Cognitive Science*, 29(2), 261–290.
doi:10.1207/s15516709cog0000_7
- Heathcote, A. & Love, J. (2012). Linear deterministic accumulator models of simple choice. *Frontiers in Psychology*, 3. doi:10.3389/fpsyg.2012.00292
- Hintze, J. L. & Nelson, R. D. (1998). Violin plots: A box plot-density trace synergism. *The American Statistician*, 52(2), 181–184. doi:10.1080/00031305.1998.10480559
- Jäger, L. A., Engelmann, F., & Vasishth, S. (2015). Retrieval interference in reflexive processing: Experimental evidence from Mandarin, and computational modeling. *Frontiers in Psychology*, 6(617). doi:10.3389/fpsyg.2015.00617
- Jäger, L. A., Engelmann, F., & Vasishth, S. (2017). Similarity-based interference in sentence comprehension: Literature review and Bayesian meta-analysis. *Journal of Memory and Language*, 94, 316–339.
- Kush, D. & Phillips, C. (2014). Local anaphor licensing in an SOV language: Implications for retrieval strategies. *Frontiers in Psychology*, 5(1252).
doi:10.3389/fpsyg.2014.01252

- Lebiere, C. (1999). The dynamics of cognition: An ACT-R model of cognitive arithmetic. *Kognitionswissenschaft*, 8(1), 5–19. doi:10.1007/BF03354932
- Lebiere, C., Anderson, J. R., & Reder, L. M. (1994). Error modeling in the ACT-R production system. In *Proceedings of the Sixteenth Annual Conference of the Cognitive Science Society* (pp. 555–559). Erlbaum Hillsdale, NJ.
- Lee, M. D. (2011). How cognitive modeling can benefit from hierarchical Bayesian models. *Journal of Mathematical Psychology*, 55(1), 1–7. doi:10.1016/j.jmp.2010.08.013
- Lewandowski, D., Kurowicka, D., & Joe, H. (2009). Generating random correlation matrices based on vines and extended onion method. *Journal of multivariate analysis*, 100(9), 1989–2001. doi:10.1016/j.jmva.2009.04.008
- Lewis, R. L. & Vasishth, S. (2005). An activation-based model of sentence processing as skilled memory retrieval. *Cognitive Science*, 29(3), 375–419. doi:10.1207/s15516709cog0000_25
- Lewis, R. L., Vasishth, S., & Van Dyke, J. A. (2006). Computational principles of working memory in sentence comprehension. *Trends in Cognitive Sciences*, 10(10), 447–454. doi:10.1016/j.tics.2006.08.007
- Logačev, P. & Vasishth, S. (2016). A multiple-channel model of task-dependent ambiguity resolution in sentence comprehension. *Cognitive Science*, 40(2), 266–298. doi:10.1111/cogs.12228
- Marcus, G. F. (2013). Evolution, Memory, and the Nature of Syntactic Representation. In Johan J. Bolhuis & M. Everaert (Eds.), *Birdsong, speech, and language: Exploring the evolution of mind and brain* (Chap. 2, pp. 27–44). MIT press.
- McElree, B. (1993). The locus of lexical preference effects in sentence comprehension: A time-course analysis. *Journal of Memory and Language*, 32(4), 536–571. doi:10.1006/jmla.1993.1028

- McElree, B. (2000). Sentence comprehension is mediated by content-addressable memory structures. *Journal of Psycholinguistic Research*, 29(2), 111–123.
doi:10.1023/A:1005184709695
- McElree, B. (2006). Accessing recent events. *Psychology of Learning and Motivation* Volume 46, 155–200. doi:10.1016/S0079-7421(06)46005-9
- McElree, B. & Doshier, B. A. (1989). Serial position and set size in short-term memory: The time course of recognition. *Journal of Experimental Psychology: General*, 118(4), 346–373. doi:10.1037/0096-3445.118.4.346
- McElree, B., Foraker, S., & Dyer, L. (2003). Memory structures that subserve sentence comprehension. *Journal of Memory and Language*, 48(1), 67–91.
doi:10.1016/s0749-596x(02)00515-6
- Nairne, J. S. (1990). A feature model of immediate memory. *Memory & Cognition*, 18(3), 251–269. doi:10.3758/BF03213879
- Nicenboim, B., Engelmann, F., Suckow, K., & Vasishth, S. (Submitted). Number interference in German: Evidence for cue-based retrieval. Open Science Framework.
doi:10.17605/OSF.IO/MMR7S
- Nicenboim, B., Logačev, P., Gattei, C., & Vasishth, S. (2016). When high-capacity readers slow down and low-capacity readers speed up: Working memory and locality effects. *Frontiers in Psychology*, 7(280). doi:10.3389/fpsyg.2016.00280
- Nicenboim, B. & Vasishth, S. (2016). Statistical methods for linguistic research: Foundational Ideas - Part II. *Language and Linguistics Compass*, 10(11), 591–613.
doi:10.1111/lnc3.12207. eprint: <https://arxiv.org/abs/1602.00245>
- Nicenboim, B., Vasishth, S., Gattei, C., Sigman, M., & Kliegl, R. (2015). Working memory differences in long-distance dependency resolution. *Frontiers in Psychology*, 1.
doi:10.3389/fpsyg.2015.00312
- Oberauer, K. (2013). The focus of attention in working memory—from metaphors to mechanisms. *Frontiers in Human Neuroscience*, 7, 673. doi:10.3389/fnhum.2013.00673

- Oberauer, K. & Kliegl, R. (2006). A formal model of capacity limits in working memory. *Journal of Memory and Language*, 55(4), 601–626. doi:10.1016/j.jml.2006.08.009
- Papaspiliopoulos, O., Roberts, G. O., & Sköld, M. (2007). A general framework for the parametrization of hierarchical models. *Statistical Science*, 22(1), 59–73. doi:10.1214/088342307000000014
- Parker, D. & Phillips, C. (2016, June 29). Negative polarity illusions and the format of hierarchical encodings in memory. *Cognition*. in press.
- Parker, D. & Phillips, C. (2017, January 26). Reflexive attraction in comprehension is selective. *Journal of Memory and Language*, 94, 272–290.
- R Core Team. (2015). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing. Vienna, Austria.
- Raab, D. H. (1962). Division of psychology: Statistical facilitation of simple reaction times. *Transactions of the New York Academy of Sciences*, 24(5 Series II), 574–590.
- Raaijmakers, J. G. W. & Shiffrin, R. M. (1980). SAM: A theory of probabilistic search of associative memory. *The psychology of Learning and Motivation*, 14, 207–262.
- Ratcliff, R. (1978). A theory of memory retrieval. *Psychological Review*, 85(2), 59. doi:10.1037/0033-295X.85.2.59
- Ratcliff, R. & McKoon, G. (2008). The diffusion decision model: Theory and data for two-choice decision tasks. *Neural Computation*, 20(4), 873–922. doi:10.1162/neco.2008.12-06-420
- Ratcliff, R. & Rouder, J. N. (1998). Modeling Response Times for Two-Choice Decisions. *Psychological Science*, 9(5), 347–356. doi:10.1111/1467-9280.00067
- Ratcliff, R., Smith, P. L., Brown, S. D., & McKoon, G. (2016). Diffusion decision model: Current issues and history. *Trends in Cognitive Sciences*, 20(4), 260–281. doi:10.1016/j.tics.2016.01.007
- Reed, A. V. (1973). Speed-accuracy trade-off in recognition memory. *Science*, 181(4099), 574–576. doi:10.1126/science.181.4099.574

- Rouder, J. N. (2005). Are unshifted distributional models appropriate for response time? *Psychometrika*, *70*(2), 377–381. doi:10.1007/s11336-005-1297-7
- Rouder, J. N., Province, J. M., Morey, R. D., Gomez, P., & Heathcote, A. (2014). The lognormal race: A cognitive-process model of choice and latency with desirable psychometric properties. *Psychometrika*, *80*(2), 491–513. doi:10.1007/s11336-013-9396-3
- Shiffrin, R. M., Lee, M., Kim, W., & Wagenmakers, E.-J. (2008). A survey of model evaluation approaches with a tutorial on hierarchical Bayesian methods. *Cognitive Science*, *32*(8), 1248–1284. doi:10.1080/03640210802414826
- Sorensen, T., Hohenstein, S., & Vasishth, S. (2016). Bayesian linear mixed models using Stan: A tutorial for psychologists, linguists, and cognitive scientists. *Quantitative Methods for Psychology*. In Press.
- Spiegelhalter, D. J., Best, N. G., Carlin, B. P., & Van Der Linde, A. (2002). Bayesian measures of model complexity and fit. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, *64*(4), 583–639. doi:10.1111/1467-9868.00353
- Stan Development Team. (2016a). Stan: A C++ library for probability and sampling, version 2.9.0.
- Stan Development Team. (2016b). *Stan modeling language users guide and reference manual, version 2.9*.
- Stan Development Team. (2017). Rstan: the R interface to Stan, version 2.9.0.
- Summerfield, C. & Tsetsos, K. (2015). Do humans make good decisions? *Trends in Cognitive Sciences*, *19*(1), 27–34. doi:10.1016/j.tics.2014.11.005
- Swets, B., Desmet, T., Clifton, C. J., & Ferreira, F. (2008). Underspecification of syntactic ambiguities: evidence from self-paced reading. *Memory & Cognition*, *36*(1), 201–216. doi:10.3758/MC.36.1.201

- Traxler, M. J., Pickering, M. J., & Clifton, C. J. (1998). Adjunct attachment is not a form of lexical ambiguity resolution. *Journal of Memory and Language*, 39(4), 558–592. doi:http://dx.doi.org/10.1006/jmla.1998.2600
- Usher, M. & McClelland, J. L. (2001). The time course of perceptual choice: The leaky, competing accumulator model. *Psychological Review*, 108(3), 550. doi:10.1037/0033-295X.108.3.550
- Van Dyke, J. A. (2007). Interference effects from grammatically unavailable constituents during sentence processing. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 33(2), 407. doi:10.1037/0278-7393.33.2.407
- Van Dyke, J. A. & Lewis, R. L. (2003). Distinguishing effects of structure and decay on attachment and repair: A cue-based parsing account of recovery from misanalyzed ambiguities. *Journal of Memory and Language*, 49(3), 285–316. doi:10.1016/S0749-596X(03)00081-0
- Van Dyke, J. A. & McElree, B. (2006). Retrieval interference in sentence comprehension. *Journal of Memory and Language*, 55(2), 157–166. doi:10.1016/j.jml.2006.03.007
- Van Dyke, J. A. & McElree, B. (2011). Cue-dependent interference in comprehension. *Journal of Memory and Language*, 65(3), 247–263. doi:10.1016/j.jml.2011.05.002
- van Gompel, R. P. G., Pickering, M. J., & Traxler, M. J. (2000). Unrestricted race: A new model of syntactic ambiguity resolution. *Reading as a perceptual process*, 621–648. doi:/10.1016/B978-008043642-5/50029-2
- Van Rijn, H. & Anderson, J. R. (2003). Modeling lexical decision as ordinary retrieval. In *In Proceedings of the International Conference on Cognitive Modeling (ICCM)* (pp. 207–212).
- van Maanen, L., van Rijn, H., & Taatgen, N. (2011). RACE/A: An architectural account of the interactions between learning, task control, and retrieval dynamics. *Cognitive Science*, 36(1), 62–101. doi:10.1111/j.1551-6709.2011.01213.x

- Vasishth, S., Brüssow, S., Lewis, R. L., & Drenhaus, H. (2008). Processing polarity: How the ungrammatical intrudes on the grammatical. *Cognitive Science*, 32(4), 685–712. doi:10.1080/03640210802066865
- Vasishth, S. & Lewis, R. L. (2006). Argument-head distance and processing complexity: Explaining both locality and antilocality effects. *Language*, 82(4), 767–794. doi:10.1353/lan.2006.0236
- Vehtari, A., Gelman, A., & Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*, 27(5), 1413–1432. doi:10.1007/s11222-016-9696-4
- Vehtari, A. & Ojanen, J. (2012). A survey of Bayesian predictive methods for model assessment, selection and comparison. *Statistics Surveys*, 6, 142–228. doi:10.1214/12-ss102
- Vickers, D. (1970). Evidence for an accumulator model of psychophysical discrimination. *Ergonomics*, 13(1), 37–58. doi:10.1080/00140137008931117
- Wagenmakers, E.-J., Ratcliff, R., Gomez, P., & McKoon, G. (2008). A diffusion model account of criterion shifts in the lexical decision task. *Journal of Memory and Language*, 58(1), 140–159. doi:10.1016/j.jml.2007.04.006
- Wagers, M. W., Lau, E. F., & Phillips, C. (2009). Agreement attraction in comprehension: representations and processes. *Journal of Memory and Language*, 61(2), 206–237. doi:10.1016/j.jml.2009.04.002
- Wickelgren, W. A. (1977). Speed-accuracy tradeoff and information processing dynamics. *Acta psychologica*, 41(1), 67–85. doi:10.1016/0001-6918(77)90012-9
- Zandbelt, B., Purcell, B. A., Palmeri, T. J., Logan, G. D., & Schall, J. D. (2014). Response times from ensembles of accumulators. *Proceedings of the National Academy of Sciences*, 111(7), 2848–2853. doi:10.1073/pnas.1310577111

Appendix A

Implementation of the activation-based model in Stan

The Stan code (shown in Listing 1) was fit to a Latin-squared design, where only the sentences of the original experiment (Nicenboim et al., Submitted) with questions that queried the subject of the embedded verb was kept, and it used a non-centered parameterization to improve convergence (for details see: Papaspiliopoulos, Roberts, & Sköld, 2007; Stan Development Team, 2016b) in Stan (Stan Development Team, 2016a). However, to improve clarity, we ignore that each participant did not respond to each experimental item, and we assume a centered parametrization in the equations below.

Let $i = 1, \dots, N_{subj}$, $j = 1, \dots, N_{items}$, and $c = 1, \dots, N_{choices}$ index participants, items, and choices in the multiple-choice questions (1 is the correct response, the target of the retrieval, 2 and 3 are incorrect responses, the competitors, and 4 is the option “I don’t know”, which represents a failed retrieval) respectively. Let $w_{i,j}$, and $RT_{i,j}$ denote the response selected and the reading times at the auxiliary verb (*hatte*) for subject i to the item j . Then we assume that reading times have the following distribution:

$$RT_{i,j} \sim \psi_i + \text{lognormal}(b - \alpha_{i,j,c=w}, \sigma) \quad (23)$$

where ψ_i is a by-subject shift, b is an arbitrary threshold (set to 10), and $\alpha_{i,j,c=w}$ represents the rate of accumulation of the “winner” accumulator. The rest of the accumulators that did not win the race must have been slower in that specific trial. From this it follows that the accumulators that lost the race have a potential $RT_{i,j,\forall c,c \neq w}$ which is larger than the observed value $RT_{i,j}$.

If all the answers are selected at least once (and if not, we can safely remove the accumulator since its rate of accumulation is so low that it never wins), the race turns into a problem of censored data, where the reading times, $RT_{i,j,\forall c,c \neq w}$, below a lower bound, $RT_{i,j,c=w}$, never occur. In order to calculate the posterior of the rate of accumulation, α , of all the accumulators, we cannot ignore the censored data (Gelman, Carlin, et al., 2014, pp.

224-227). However, it is not necessary to impute values, and the values can be integrated out (Stan Development Team, 2016b, pp. 107–110; Gelman, Carlin, et al., 2014, pp. 224-227). Each censored data point has a probability of

$$Pr[RT_{n,\forall c,c \neq w} > RT_{n,c=w}] = \int_{RT_{n,c=w}}^{\infty} \text{lognormal}(RT_{n,\forall c,c \neq w} - \psi | b - \alpha_{\forall c,c \neq w}, \sigma) \cdot dRT_{n,\forall c,c \neq w} \quad (24)$$

$$= 1 - \Phi\left(\frac{\log(RT_{n,c=w} - \psi) - (b - \alpha_{\forall c,c \neq w})}{\sigma}\right) \quad (25)$$

where $\Phi()$ is the cumulative distribution function of the standard normal distribution. Since the shifts of the distribution, ψ_i , must be positive, to ensure convergence of the model we exponentiate a term that is associated with a general shift of the whole reading times distribution, ψ' , and a term that represents the by-participants adjustment, ψ'_i :

$$\psi_i = \exp(\psi' + \psi'_i) \quad (26)$$

with the following priors for the by-participant component:

$$\psi'_i \sim \text{normal}(0, \tau_\psi) \quad (27)$$

$$\tau_\psi \sim \text{normal}(0, 10); \text{ with } \tau_\psi > 0 \quad (28)$$

In addition, each ψ_i must be smaller than the shortest reading time of each participant i (recall that the shift is the lower bound of the distribution). We satisfied the constraint on the upper bound with the following prior on the general shift. We show in Eq. (29) the parameter ψ' back-transformed from unit-scale; notice that the standard deviation of the normal prior is in log-scale (multiplied by ten to have a weaker prior), because this parameter is exponentiated in (28) and forms the shift of the distribution.

$$\psi' \sim \text{normal}(0, 10 \cdot \log(\text{mean}(RT))); \text{ with } \psi' < U \quad (29)$$

where a normal distribution is truncated on the upper limit, U , which is the smallest difference between $\log(RT)$ and ψ'_i .

We assume that the rates of accumulation depend on the experimental condition (high or low interference) and that the rates may be affected by participants and by items. We can express this in matrix notation for each accumulator as follows:

$$\alpha_c = \mathbf{X}\boldsymbol{\beta}_c + \mathbf{X}\mathbf{u}_c + \mathbf{X}\mathbf{v}_c \quad (30)$$

Here \mathbf{X} is the $N_{obs} \times N_{pars}$ model matrix, with the number of parameters (so-called fixed effects), N_{pars} , being two: intercept and condition. Each $\boldsymbol{\beta}_c$ is a vector of length N_{pars} with the estimates of the fixed-effect parameters for the accumulator associated with the choice c . Each \mathbf{u}_c and \mathbf{v}_c are the by-participants and by-item adjustments to the fixed effects estimates (so-called random-effects) for the accumulator c . We used weakly informative priors for all the parameters (some estimates were reparametrized in the Stan implementation, see Listing 1 for details).

Here, $\beta_{0,c}$ are the intercepts of the fixed effects for choice c , $\log(\text{mean}(RT_c))$ is the logarithm of the mean of the reading times when option c was selected, and $\beta_{1,c}$ represents the slopes of the fixed effects (i.e., the effect of interference).

All the random-effects, \mathbf{u}_c , and \mathbf{v}_c , were assumed to be sampled from two multivariate normal distributions with means of zero. The prior of the standard deviations of the random effects was $\text{normal}(0, 10)$. We placed lkj priors on the random effects correlation matrices with shape parameter $\eta = 2$ (see Lewandowski, Kurowicka, & Joe, 2009; Sorensen, Hohenstein, & Vasishth, 2016).

Listing 1: Stan code for the activation-based model


```

1 functions {
2   vector mean_of_X_by_Y(vector X, int[] Y){
3     int N_rows;
4     int N_groups;
5     N_rows = num_elements(X);
6     if(N_rows!= num_elements(Y))
7       reject("X and Y don't have the same length");
8     N_groups = max(Y);
9     { # matrix with a column for each group of Y,
10      #and 1 if Y belong to the group:
11      matrix[N_rows, N_groups] matrix_1;
12      for(r in 1:N_rows)
13        for(g in 1:N_groups)
14          matrix_1[r,g] = (g == Y[r]);
15      return((X' * matrix_1) // sum of Xs per group
16             // divided by matrix_1^T * matrix_1 (number of times each group appears)
17             :
18             / crossprod(matrix_1))');
19    }
20   real psi_max(vector u_psi, int[] subj, vector RT) {
21     // This function ensures that psi is correctly truncated on the correct
22     // upper limit, U in Eq. (29)
23     real psi_max;
24     psi_max = positive_infinity();
25     for (i in 1:num_elements(RT))
26       psi_max = fmin(psi_max, log(RT[i]) - u_psi[subj[i]]);
27     return (psi_max);
28   }
29   real race(int winner, real RT, row_vector alpha, real b, real sigma, real
30     psi)
31   {
32     real log_lik;
33     int N_choices;
34     real rate;
35     real shifted_RT;
36
37     N_choices = cols(alpha);
38     shifted_RT = RT - psi;
39
40     log_lik = 0;
41     for(c in 1:N_choices){
42       if(c == winner){
43         // loglik due to distribution of observed RTs; Eq. (23)
44         log_lik = log_lik + lognormal_lpdf(shifted_RT|b - alpha[c], sigma);
45       } else {
46         // loglik due to censored RTs of loser candidates; Eq. (24)
47         log_lik = log_lik + lognormal_lccdf(shifted_RT|b - alpha[c], sigma);
48       }
49     }
50     return(log_lik);
51   }

```

```

52 data {
53   int<lower = 0> N_obs;
54   int<lower=1> N_subj;
55   int<lower=1> N_item;
56   int<lower = 1> N_choices;
57   int<lower = 1, upper = N_choices> winner[N_obs]; // response selected
58   vector<lower = 0>[N_obs] RT;
59   // for fixed effects
60   int N_coef; //intercept + predictors
61   matrix[N_obs,N_coef] x;
62   // for random effects by subject:
63   int<lower=0> N_coef_u;
64   int<lower=1> subj[N_obs];
65   matrix[N_obs,N_coef_u] x_u;
66   // for random effects by items:
67   int<lower=0> N_coef_w;
68   int<lower=1> item[N_obs];
69   matrix[N_obs,N_coef_w] x_w;
70
71 }
72 transformed data {
73   real b;
74   real min_RT;
75   real logmean_RT;
76   row_vector[N_choices] logmean_RT_w;
77   matrix[N_obs,N_coef-1] x_betas;
78   int N_tau_u;
79   int N_tau_w;
80
81   b = 10;
82   min_RT = min(RT);
83   logmean_RT = log(mean(RT));
84   logmean_RT_w = (log(mean_of_X_by_Y(RT, winner)))';
85   x_betas = x[,2:N_coef]; // x without intercept
86   N_tau_u = N_coef_u * N_choices;
87   N_tau_w = N_coef_w * N_choices;
88 }
89 parameters{
90   row_vector[N_choices] beta_0raw;
91   vector<lower = 0> [N_tau_u] tau_u;
92   cholesky_factor_corr[N_tau_u] L_u;
93   matrix[N_tau_u, N_subj] z_u;
94   vector<lower = 0> [N_tau_w] tau_w;
95   cholesky_factor_corr[N_tau_w] L_w;
96   matrix[N_tau_w, N_item] z_w;
97   real<lower = 0> sigma;
98   vector[N_subj] u_psi;
99   real<lower = 0> tau_psi;
100   real<upper = psi_max(u_psi, subj, RT) / logmean_RT> psi_p_raw;
101   matrix[N_coef-1,N_choices] beta;
102 }
103 transformed parameters {
104   real psi_p;
105   vector[N_obs] psi;

```

```

106 matrix[N_coef_u, N_choices] u[N_subj];
107 matrix[N_coef_w, N_choices] w[N_item];
108 matrix[N_obs, N_choices] alpha;
109 row_vector[N_choices] beta_0;
110
111 {
112   matrix[N_tau_u, N_subj] u_long;
113   matrix[N_tau_w, N_item] w_long;
114   // Optimization through Cholesky Fact:
115   u_long = diag_pre_multiply(tau_u, L_u) //matrix[N_choices, N_choices]
116           * z_u;
117   w_long = diag_pre_multiply(tau_w, L_w) //matrix[N_choices, N_choices]
118           * z_w;
119   for (i in 1:N_subj)
120     u[i] = to_matrix(u_long[, i], N_coef_u, N_choices, 0);
121   for (j in 1:N_item)
122     w[j] = to_matrix(w_long[, j], N_coef_w, N_choices, 0);
123 }
124 // Unit-scaling:
125 beta_0 = b - beta_0raw .* logmean_RT_w;
126 psi_p = psi_p_raw * logmean_RT;
127 // Fixed effects in alpha; Eq. (30):
128 alpha = rep_matrix(beta_0, N_obs) + x_betas * beta;
129 for (n in 1:N_obs) {
130   // Shift adjustment for every observation; Eq. (30):
131   psi[n] = exp(psi_p + u_psi[subj[n]]);
132   // Add random effects
133   for (uu in 1:N_coef_u)
134     alpha[n] = alpha[n] + x_u[n, uu] * u[subj[n], uu];
135   for (ww in 1:N_coef_w)
136     alpha[n] = alpha[n] + x_w[n, ww] * w[item[n], ww];
137 }
138 }
139 model {
140   vector[N_obs] log_lik;
141   beta_0raw ~ normal(0, 10);
142   to_vector(beta) ~ normal(0, 10);
143   tau_u ~ normal(0, 10);
144   L_u ~ lkj_corr_cholesky(2.0);
145   to_vector(z_u) ~ normal(0, 10);
146   tau_w ~ normal(0, 10);
147   L_w ~ lkj_corr_cholesky(2.0);
148   to_vector(z_w) ~ normal(0, 1);
149   sigma ~ normal(0, 10);
150   psi_p_raw ~ normal(0, 10);
151   u_psi ~ normal(0, tau_psi); // Eq. (27)
152   tau_psi ~ normal(0, 10); // Eq. (28)
153   for (n in 1:N_obs)
154     log_lik[n] = race(winner[n], RT[n], alpha[n], b, sigma, psi[n]);
155   target += log_lik;
156 }

```

Appendix B

Implementation of the direct-access model in Stan

The code (shown in Listing 2) was fit to the same data as the activation-based model. As before, to improve clarity, we ignore that each subject did not respond to each item and we assume a centered parametrization.

Let $i = 1, \dots, N_{subj}$, $j = 1, \dots, N_{items}$, and $c = 1, \dots, N_{choices}$ index participants, items, and choices respectively, where choice 1 is the correct response and choice $N_{choices}$ (which maps to 4) is the response associated with a retrieval failure. Let $w_{i,j}$, and $RT_{i,j}$ denote the response selected and the reading times at the auxiliary verb (*hatte*) for subject i to the item j .

We implemented the assumptions of the direct-access model, by letting w have a discrete distribution that follows a one-inflated categorical model, where additional probability mass is added to the outcome 1 (correct response) due to backtracking with probability θ_b as follows:

$$P(w_n = 1 | \boldsymbol{\theta}_{i,j}, \theta_b) = \theta_{1_{i,j}} + (1 - \theta_{1_{i,j}}) \cdot \theta_b \quad (31)$$

$$P(w_n = s | \boldsymbol{\theta}_{i,j}, \theta_b) = \theta_{s_{i,j}} \cdot (1 - \theta_b), \text{ with } s > 1 \quad (32)$$

where $\boldsymbol{\theta}$ is a vector of $N_{choices}$ rows that represents the probability of each option.

If the answer given is wrong, we assume that there is no backtracking and then reading times are distributed in the following way:

$$RT_{i,j,\forall w, w>1} \sim \psi_i + \text{lognormal}(T_{da,i,j}, \sigma) \quad (33)$$

where ψ_i is a by-subject shift, T_{da} represents the time needed for the direct access or failure together with extra processes

If the answer given is right, reading times are assumed to have a mixture distribution. This is so because there are two “paths” to reach a correct response (see Figure 4): (i) The chunk that is retrieved is the correct one (at the first try), and this means that there is direct access and reading times should belong to a distribution similar to the previous one as shown in Eq. (33); or (ii) an incorrect chunk (or no chunk) is retrieved but is backtracked, and this means that reading times should belong to a distribution with a larger location than $T_{da,i,j}$, namely, $T_{da,i,j} + t_{b,i,j}$. Thus RTs should be distributed in the following way:

$$RT_{i,j,w=1} \sim \psi_i + \begin{cases} \text{lognormal}(T_{da,i,j}, \sigma); & \text{if } y = 1 | \text{Categorical}(y | \boldsymbol{\theta}_{i,j}) \\ \text{lognormal}(T_{da,i,j} + t_{b,i,j}, \sigma); & \text{if } y \neq 1 \text{ and } z = 1 | \text{Categorical}(y | \boldsymbol{\theta}_{i,j}) \text{ and } \text{Bernoulli}(z | \theta_b) \end{cases} \quad (34)$$

where, from Eq. (31), the first component of the mixture defined in (34) is the probability of a correct retrieval at the first attempt conditional on obtaining *one* as a

response ($w_n = 1$), and occurs with probability:

$$\frac{\theta_{1,i,j}}{\theta_{1,i,j} + (1 - \theta_{1,i,j}) \cdot \theta_b} \quad (35)$$

and the second component of the mixture (34) represents the probability of backtracking when there is an error, this is formulated as the probability of an incorrect retrieval in the first attempt conditional on obtaining *one* as a response ($w_n = 1$):

$$\frac{1 - \theta_{1,i,j}}{\theta_{1,i,j} + (1 - \theta_{1,i,j}) \cdot \theta_b} \quad (36)$$

Furthermore, the categorical model was fit including a hierarchical structure in θ' , where θ' is a vector with $N_{choices}$ rows with its last row set to zero, so that $softmax(\theta') = \theta$.⁶ This way we ensure the identifiability of $Categorical(softmax(\theta'))$.

We assume that the probability of each choice depend on the experimental condition (high or low interference) and that the probabilities may be affected by participants and by items. In matrix notation, the first $N_{choices} - 1$ rows of θ' are structured as the activations in the activation-based model:

$$\theta'_c = \mathbf{X}\beta_c + \mathbf{X}\mathbf{u}_c + \mathbf{X}\mathbf{v}_c \quad (37)$$

As for the activation-based model, we used weakly informative priors for all the estimates. The priors for the fixed effects were set with the added constraint that $\beta_{0,1}$, the intercept of the probability of the correct choice (the first choice) in logit-space, was constrained to be larger than $\beta_{0,2..3}$, the intercept associated with the incorrect responses, and zero (which is the value associated with the last choice):

$$\beta_{0,2..3} \sim normal(0, 10) \quad (38)$$

$$\beta_{0,1} \sim normal(0, 10) + max(\beta_{0,2..3}, 0) \quad (39)$$

$$\theta_b \sim beta(1, 1) \quad (40)$$

In addition, we assumed a hierarchical structure to the parameters associated with latencies:

$$T_{da,i,j} = \beta_{0,Tda} + u_{Tda,i} + v_{Tda,j} \quad (41)$$

$$T_{b,i,j} = \beta_{0,Tb} + u_{Tb,i} + v_{Tb,j} \quad (42)$$

with the following priors on the intercepts:

$$\beta_{0,Tda} \sim normal(0, 10 \cdot \log(mean(RT))) \quad (43)$$

$$\beta_{0,Tb} \sim normal(0, 10) \quad (44)$$

⁶The softmax function is defined as in Stan Development Team (2016b) by $softmax(y) = \frac{\exp(y)}{\sum_{k=1}^K \exp(y_k)}$

All the random-effects, \mathbf{u}_c , u_{tda} , u_{tb} , \mathbf{v}_c , v_{tda} , v_{tb} were assumed to be sampled from four multivariate normal distributions with means of zero: (i) for the subject adjustment on probabilities of retrieval, (ii) for a similar adjustment for items, (iii) for the subject adjustment on latencies, and (iv) for a similar adjustment for items. As before we placed lkj priors on the random effects correlation matrices with shape parameter $\eta = 2$.

Listing 2: Stan code for the direct-access model

```

1 functions {
2   real psi_max(vector u_psi, int[] subj, vector RT) {
3     real psi_max;
4     psi_max = positive_infinity();
5     for (i in 1:num_elements(RT))
6       psi_max = fmin(psi_max, log(RT[i]) - u_psi[subj[i]]);
7     return (psi_max);
8   }
9   real da(int winner, real RT, row_vector thetap, real theta_b, real T_da,
10          real T_b, real sigma, real psi){
11     vector[num_elements(thetap)] thetapT;
12     // theta = softmax(thetapT)
13
14     real log_P_w1; // log(P(w = 1 | theta, theta_b)); Eq. (31)
15     real log_P_da_gw1; // log(Prob of direct access given winner = 1); Eq. (35)
16     real log_theta_b_gw1; // log(Prob of backtracking given winner = 1); Eq.
17                          (36)
18     thetapT = thetap';
19
20     // Eq. (31) in log:
21     log_P_w1 = log_sum_exp(categorical_logit_lpmf(1 | thetapT),
22                           log(theta_b) + loglm_exp(categorical_logit_lpmf(1|thetapT)));
23     // Eq. (35) in log:
24     log_P_da_gw1 = categorical_logit_lpmf(1 | thetapT) - log_P_w1;
25     // Eq. (36) in log:
26     log_theta_b_gw1 = log(theta_b) +
27                      loglm_exp(categorical_logit_lpmf(1 | thetapT)) - log_P_w1;
28     if(winner==1) {
29       return (log_P_w1 + // Increment on likelihood due to winner=1
30               // Increment on likelihood due to RT:
31               log_sum_exp(log_P_da_gw1 + lognormal_lpdf(RT - psi | T_da, sigma),
32                           log_theta_b_gw1 + lognormal_lpdf(RT - psi | T_da + T_b, sigma) ));
33     } else {
34       return (loglm(theta_b) + //
35               categorical_logit_lpmf(winner | thetapT) + // Eq. (32)
36               // Increment on likelihood due to RTs; Eq. (33):
37               lognormal_lpdf(RT - psi | T_da, sigma));
38     }
39   }
40 }
41 data {
42   int<lower = 0> N_obs;
43   int<lower = 1> N_choices;
44   int<lower = 1, upper = N_choices> winner[N_obs];

```

```

45   vector<lower = 0>[N_obs] RT;
46   int N_coef;
47   matrix[N_obs,N_coef] x;
48   int<lower = 0> N_coef_u;
49   int<lower = 1> subj[N_obs];
50   int<lower = 1> N_subj;
51   matrix[N_obs, N_coef_u] x_u;
52   int<lower = 0> N_coef_w;
53   int<lower = 1> item[N_obs];
54   int<lower = 1> N_item;
55   matrix[N_obs, N_coef_w] x_w;
56 }
57 transformed data {
58   real<lower = 0> min_RT;
59   real logmean_RT;
60   matrix[N_obs, N_coef - 1] x_betas;
61   int N_tau_u;
62   int N_tau_w;
63
64   min_RT = min(RT);
65   logmean_RT = log(mean(RT));
66   x_betas = x[, 2:N_coef]; // intercept removed
67   N_tau_u = N_coef_u * (N_choices - 1);
68   N_tau_w = N_coef_w * (N_choices - 1);
69 }
70 parameters{
71   real<lower = 0> sigma;
72   real<lower = 0> beta_0_Tdaraw;
73   real<lower = 0> beta_0_Tb;
74   row_vector[N_choices-2] thetap_incorrect;
75   real<lower = 0> thetap_added;
76   matrix[N_coef-1,N_choices-1] beta;
77   vector<lower = 0> [N_tau_u] tau_u;
78   cholesky_factor_corr[N_tau_u] L_u;
79   matrix[N_tau_u, N_subj] z_u;
80   vector<lower = 0> [2] tau_u_RT;
81   cholesky_factor_corr[2] L_u_RT;
82   matrix[2, N_subj] z_u_RT;
83   vector<lower = 0> [N_tau_w] tau_w;
84   cholesky_factor_corr[N_tau_w] L_w;
85   matrix[N_tau_w, N_item] z_w;
86   vector<lower = 0> [2] tau_w_RT;
87   cholesky_factor_corr[2] L_w_RT;
88   matrix[2, N_item] z_w_RT;
89   real<lower = 0, upper = 1> theta_b;
90   vector[N_subj] u_psi;
91   real<lower = 0> tau_psi;
92   real<upper = psi_max(u_psi, subj, RT) / logmean_RT> psi_p_raw;
93 }
94 transformed parameters{
95   real<lower=0> beta_0_Tda;
96   vector[N_obs] T_da;
97   vector[N_obs] T_b;
98   vector[N_obs] psi;

```

```

99  matrix[2, N_subj] u_RT;
100 matrix[N_coef_u, N_choices - 1] u[N_subj];
101 matrix[2, N_item] w_RT;
102 matrix[N_coef_w, N_choices - 1] w[N_item];
103 real psi_p;
104 matrix[N_obs, N_choices] thetap;
105 row_vector[N_choices-1] beta_0;
106
107 {
108     matrix[N_tau_u, N_subj] u_long;
109     matrix[N_tau_w, N_item] w_long;
110     // Optimization through Cholesky Fact:
111     u_RT = diag_pre_multiply(tau_u_RT, L_u_RT) * z_u_RT;
112     w_RT = diag_pre_multiply(tau_w_RT, L_w_RT) * z_w_RT;
113     u_long = diag_pre_multiply(tau_u, L_u) //matrix[N_tau_u, N_tau_u]
114             * z_u;
115     w_long = diag_pre_multiply(tau_w, L_w) //matrix[N_tau_w, N_tau_w]
116             * z_w;
117     for (i in 1:N_subj)
118         u[i] = to_matrix(u_long[, i], N_coef_u, (N_choices-1), 0);
119     for (j in 1:N_item)
120         w[j] = to_matrix(w_long[, j], N_coef_w, (N_choices-1), 0);
121 }
122
123 beta_0[1] = thetap_added + fmax(max(thetap_incorrect), 0); // Eq. (39)
124 beta_0[2:] = thetap_incorrect; //Eq. (38)
125 beta_0_Tda = beta_0_Tdaraw * logmean_RT; // Eq. (43)
126 psi_p = psi_p_raw * logmean_RT;
127 // Adds so called fixed effects; first summand of Eq. (37)
128 thetap[, 1:N_choices-1] = rep_matrix(beta_0, N_obs) + x_betas * beta ;
129 thetap[, N_choices] = rep_vector(0, N_obs);
130
131 for (n in 1:N_obs) {
132     psi[n] = exp(psi_p + u_psi[subj[n]]);
133     T_da[n] = beta_0_Tda + u_RT[1, subj[n]] + w_RT[1, item[n]]; // Eq. (41)
134     T_b[n] = beta_0_Tb + u_RT[2, subj[n]] + w_RT[2, item[n]]; // Eq. (42)
135     //adds so called random effects; second two summands of Eq. (37)
136     for (uu in 1:N_coef_u)
137         thetap[n, 1:N_choices - 1] = thetap[n, 1:N_choices - 1] + x_u[n, uu] *
138                                     u[subj[n], uu];
139     for (ww in 1:N_coef_w)
140         thetap[n, 1:N_choices - 1] = thetap[n, 1:N_choices - 1] + x_w[n, ww] *
141                                     w[item[n], ww];
142 }
143 }
144 model {
145     vector[N_obs] log_lik;
146     sigma ~ normal(0, 10);
147     thetap_added ~ normal(0, 10); // First summand of Eq. (39)
148     thetap_incorrect ~ normal(0, 10); //Eq. (38)
149     to_vector(beta) ~ normal(0, 1);
150     psi_p_raw ~ normal(0, 10);
151     u_psi ~ normal(0, tau_psi); // Eq. (27)
152     tau_psi ~ normal(0, 10); // Eq. (28)

```



```

153   to_vector(z_u_RT) ~ normal(0, 1);
154   to_vector(z_u) ~ normal(0, 1);
155   tau_u_RT ~ normal(0, 10);
156   tau_u ~ normal(0, 10);
157   L_u_RT ~ lkj_corr_cholesky(2.0);
158   L_u ~ lkj_corr_cholesky(2.0);
159   to_vector(z_w_RT) ~ normal(0, 1);
160   to_vector(z_w) ~ normal(0, 1);
161   tau_w_RT ~ normal(0, 10);
162   tau_w ~ normal(0, 10);
163   L_w_RT ~ lkj_corr_cholesky(2.0);
164   L_w ~ lkj_corr_cholesky(2.0);
165   theta_b ~ beta(1, 1); // Eq. (40)
166   beta_0_Tdaraw ~ normal(0, 10);
167   beta_0_Tb ~ normal(0, 10); // Eq. (44)
168   for (n in 1:N_obs)
169     log_lik[n] = da(winner[n], RT[n], thetap[n], theta_b, T_da[n], T_b[n],
170                    sigma,
171                    psi[n]);
171   target += log_lik;
172 }

```

Appendix C

Recovery of the parameters

We show here the ability of the models to recover the parameters generated from fake datasets. We used three fake datasets, one for testing each of the models we discussed in the paper: the activation-based model, the direct-access model, and the activation-based model with different variances. The procedure for each dataset and corresponding model is as follows:

1. We extract the point estimates of a given model.
2. We generated a dataset assuming that the given model is the *true generating process*, and the true value of the parameters are the means of the posteriors from 1.
3. We fit the model to its corresponding dataset.
4. We extract the posteriors of the estimated parameters.
5. We show graphically the discrepancy between the estimated posteriors and the true values.

Figures C1, C2, and C3 show the true values, and the mean and 95% CrI of the posterior distribution for the parameters of the activation-based model, direct-access model, and activation-based model with different variances respectively.




Figure C1. Discrepancies between estimated and true value of the parameters of the activation-based model. Black points indicate the mean of the posteriors and the black horizontal lines indicate the 95% CrIs, the stars indicate the true value of the parameters.




Figure C2. Discrepancies between estimated and true value of the parameters of the direct-access model. Black points indicate the mean of the posteriors and the black horizontal lines indicate the 95% CrIs, the stars indicate the true value of the parameters.




Figure C3. Discrepancies between estimated and true value of the parameters of the activation-based model with different variances. Black points indicate the mean of the posteriors and the black horizontal lines indicate the 95% CrIs, the stars indicate the true value of the parameters.